

Point72 MI Data - Analytics Case Study

Introduction

Shown below is the problem statement of this case study.

- You are given two types of data: 311 Service requests, and weather. The 311 Service requests dataset contains all the 311 calls from 2016 to 2018 with many different features including timestamps, location, service request descriptions, etc. Please describe the steps you take to ingest and process the data, what are the challenges and how you overcome them.
- Aggregate the 311 records in ways that best describes the characteristics of the inbound call patterns. Explain and visualize your findings. What are the implications you can infer from the patterns?
- The weather data contains the weather information from the weather stations in NYC from 2010 to 2018 in daily intervals. Describe and visualize your findings. (Please note this is NOT weather forecast)
- The final goal of this project is to predict the daily 311 inbound calls for the next 7 days
- Use your insights from previous sections to build a predictive model of your choice – explain the assumptions of your model and why you picked it.
- What features are you using in your model and, how does your data exploration process help you decide which features you are incorporating in your model?
- Evaluate the impact of weather on the call volume, find patterns and define causal impact if there's any.
- Build a reasonable model framework, explain your model results, and justify the relationships identified in the model. Not necessary to dig into complex models.

311 requests

First, we'll look at the 311 calls dataset. This dataset in its raw form is quite large (7M rows, and 43 columns). To avoid running into memory issues, and keep runtimes reasonable, we can sample a chunk of the dataset instead of reading all of it. We can increase the sample size after the data exploration phase, if a need arises - such as if an analysis we're interested in requires more data points.

To analyze the 311 Service requests dataset, we will follow a structured approach to ingest, process, and aggregate the data, focusing on understanding the inbound call patterns. Here's an outline of the steps we'll take, the potential challenges we might encounter, and how we'll overcome them, followed by an analysis and visualization of our findings. The detailed steps described below are standard steps for data analysis, like we'd find in any reference material.

Step 1: Data Ingestion

Action: Load the dataset using Pandas.\b Challenge: The file size is very large, and leads to memory issues.\b Solution: Use chunking and randomly sample data points so we have span all the timeperiods.

Step 2: Preliminary Exploration

Action: Perform an initial exploration to understand the dataset's structure, including the number of records, features, missing values, and data types.\b Challenge: Identifying relevant features and handling missing or inconsistent data.\b Solution: Use descriptive statistics and visualization tools to assess data quality, and apply imputation or removal of missing data as appropriate.

Step 3: Data Cleaning

Action: Clean the dataset by handling missing values, correcting data types (e.g., converting timestamps to datetime objects), and removing duplicates.\b Challenge: Ensuring accurate data type conversions and dealing with outliers.\b Solution: Validate conversions through sample checks and use statistical methods to identify and handle outliers.

Step 4: Data Aggregation

Action: Aggregate the data to identify patterns, such as call volume over time, most common service requests, and geographic distribution of calls.\b Challenge: Choosing the right level of aggregation to reveal meaningful patterns without oversimplification.\b Solution: Experiment with different aggregation levels (e.g., daily, monthly, by neighborhood) and metrics (e.g., count, mean) to find the most insightful views.

Step 5: Analysis and Visualization

Action: Analyze aggregated data to uncover trends, seasonal patterns, and anomalies. Visualize findings using charts and maps.\b Challenge: Making complex data understandable and visually engaging.\b Solution: Use a combination of visualization techniques, such as time series plots, bar charts, heatmaps, and geographic maps.

Step 6: Implications and Insights

Action: Interpret patterns to infer implications for city management, resource allocation, and policy making.\b Challenge: Translating data patterns into actionable insights.\b Solution: Combine data analysis with domain knowledge to provide recommendations or insights. Let's start by ingesting the dataset and performing a preliminary exploration to understand its structure. We'll look at the number of records, features, and get a sense of the data we're dealing with. Then, we'll proceed with the subsequent steps based on our initial findings.

We'll jump right into it by importing libraries. Then we'll load and examine the data to understand its structure and contents. This examination will allow us to identify the best ways to aggregate the 311 records and describe the characteristics of the inbound call patterns. Let's start by loading the data and taking a look at the first few rows.

In [1]:

```
import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import requests
from io import StringIO
import random
import plotly.express as px

font = {'size' : 18}
plt.rcParams['font', **font]
```

In [2]:

```
filename = '311-2016-2018.csv'

n = sum(1 for line in open(filename)) - 1 # number of rows in the original dataset
s = 1000000 # we want to sample 500k rows, randomly
skip = sorted(random.sample(range(1,n+1),n-s)) # the sampling step

requests = pd.read_csv(filename, skiprows=skip, on_bad_lines='skip')
requests['Created Date'] = pd.to_datetime(requests['Created Date'])
requests.head()
```

Out [2]:

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Descriptor	Location Type	Incident Zip	Incident Address	...	Taxi Company Borough	Tax Pick Up Location
0	38237855	2018-01-17 22:42:00	01/20/2018 10:02:00 PM	DSNY	Department of Sanitation	Request Large Bulky Item Collection	Request Large Bulky Item Collection	Sidewalk	11361.0	204-36 47 ROAD	...	NaN	NaN
1	38237869	2018-01-17 16:10:25	01/25/2018 05:01:02 PM	DOT	Department of Transportation	Sidewalk Condition	Sidewalk Violation	Sidewalk	11103.0	25-92 45 STREET	...	NaN	NaN

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Descriptor	Location Type	Incident Zip	Incident Address	...	Taxi Company Borough	Tax Pick Up Location
2	38237875	2018-01-17 20:59:39	04/09/2018 04:00:34 PM	DPR	Department of Parks and Recreation	New Tree Request	For One Address	Street	10456.0	1069 MORRIS AVENUE	...	NaN	NaN
3	38237880	2018-01-17 06:14:00	01/19/2018 10:52:00 PM	DSNY	Department of Sanitation	Request Large Bulky Item Collection	Request Large Bulky Item Collection	Sidewalk	11365.0	56-43A UTOPIA PARKWAY	...	NaN	NaN
4	38237911	2018-01-16 14:28:42	01/17/2018 10:34:00 AM	DOT	Department of Transportation	Street Condition	Pothole	Nan	11434.0	NaN	...	NaN	NaN

5 rows × 42 columns

In [3]: requests.columns

```
Out[3]: Index(['Unique Key', 'Created Date', 'Closed Date', 'Agency', 'Agency Name',
       'Complaint Type', 'Descriptor', 'Location Type', 'Incident Zip',
       'Incident Address', 'Street Name', 'Cross Street 1', 'Cross Street 2',
       'Intersection Street 1', 'Intersection Street 2', 'Address Type',
       'City', 'Landmark', 'Facility Type', 'Status', 'Due Date',
       'Resolution Description', 'Resolution Action Updated Date',
       'Community Board', 'BBL', 'Borough', 'X Coordinate (State Plane)',
       'Y Coordinate (State Plane)', 'Open Data Channel Type',
       'Park Facility Name', 'Park Borough', 'Vehicle Type',
       'Taxi Company Borough', 'Taxi Pick Up Location', 'Bridge Highway Name',
       'Bridge Highway Direction', 'Road Ramp', 'Bridge Highway Segment',
       'Latitude', 'Longitude', 'Location', 'Created Year'],
      dtype='object')
```

The dataset contains various columns, including identifiers, dates, agency details, complaint types, descriptors, location information, and more. Here's a brief overview based on the first few rows:

- Unique Key: A unique identifier for each record.
- Created Date: The date and time when the complaint was created.
- Closed Date: The date and time when the complaint was closed.
- Agency: The acronym of the agency responsible for addressing the complaint.
- Agency Name: The full name of the agency.

- Complaint Type: The type of complaint.
- Descriptor: More detailed information about the complaint.
- Location Type: The type of location where the incident was reported.
- Incident Zip: The ZIP code where the incident occurred.
- Latitude, Longitude, Location: Geographic details of the incident location.
- Created Year and Date: The year and date when the complaint was created, respectively.

Exploratory data analysis

To describe the characteristics of the inbound call patterns, we can aggregate the data in several ways:

1. Time Trends: Analyze complaint volume over time (daily, monthly, yearly).
2. Complaint Type Distribution: Identify the most common types of complaints.
3. Agency Response: Examine which agencies handle the most complaints.
4. Location Analysis: Determine areas with high complaint volumes.
5. Resolution Time: Calculate the time taken to close complaints.

For a comprehensive analysis, we'll start by exploring each of these aggregations.

Yearly, monthly trends

Let's begin with the time trends to see how complaint volumes have changed over time. We'll look at the yearly and monthly complaint volumes.

In [4]:

```
# Extract year and month for aggregation
requests['Year'] = requests['Created Date'].dt.year
requests['Month'] = requests['Created Date'].dt.month

# Yearly complaint volumes
yearly_volume = requests.groupby('Year').size()

# Monthly complaint volumes (across all years)
monthly_volume = requests.groupby('Month').size()

# Plotting
fig, ax = plt.subplots(2, 1, figsize=(12, 12))

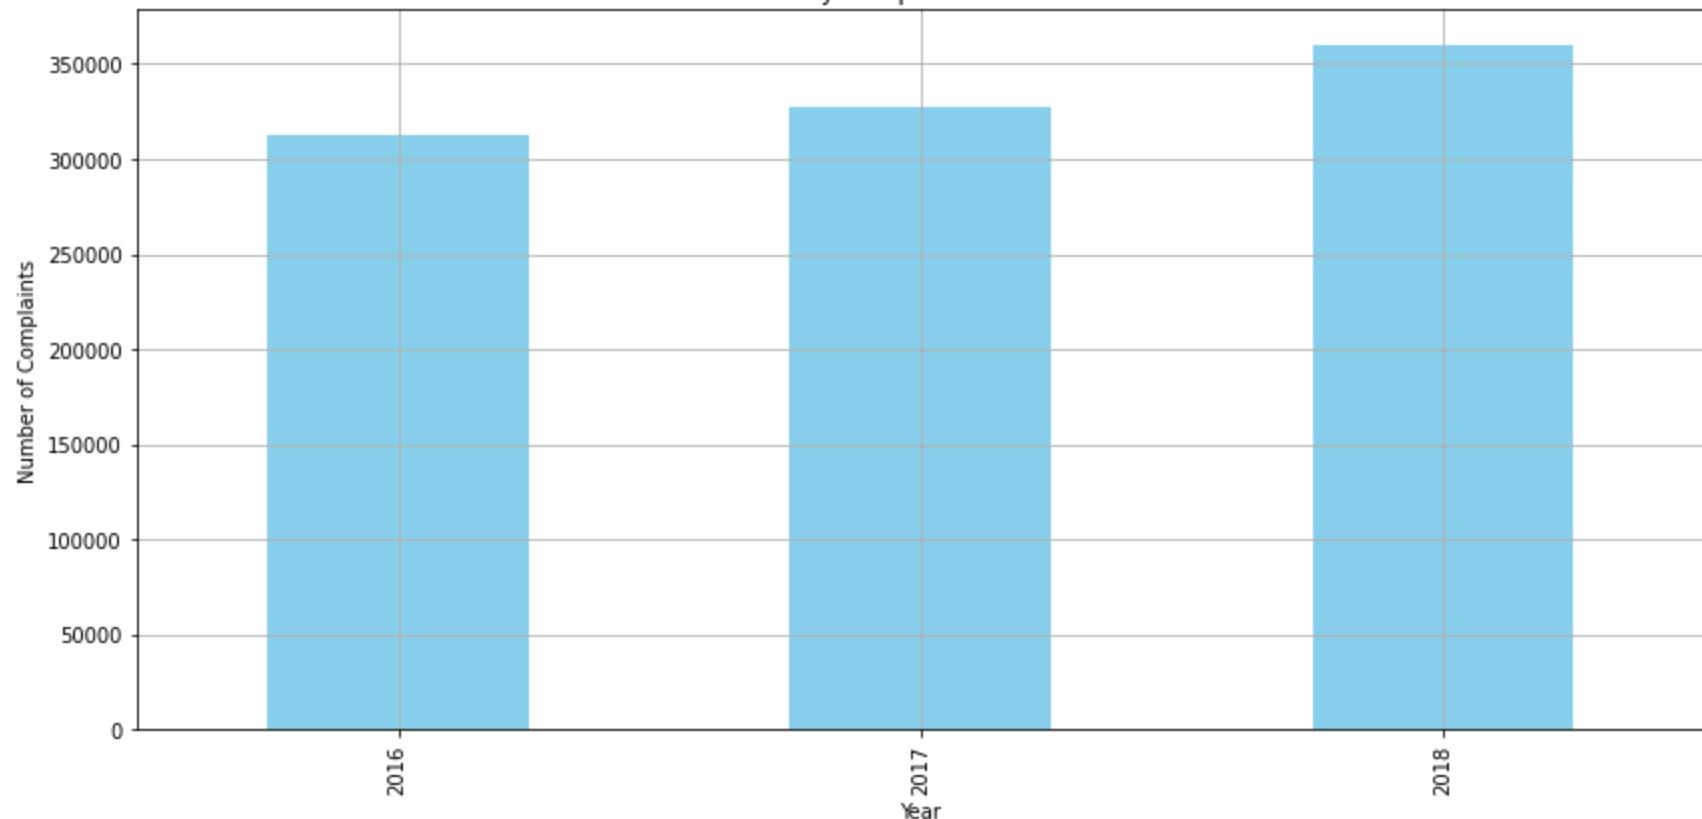
# Yearly Volumes
yearly_volume.plot(kind='bar', ax=ax[0], color='skyblue')
```

```
ax[0].set_title('Yearly Complaint Volumes')
ax[0].set_xlabel('Year')
ax[0].set_ylabel('Number of Complaints')
ax[0].grid()

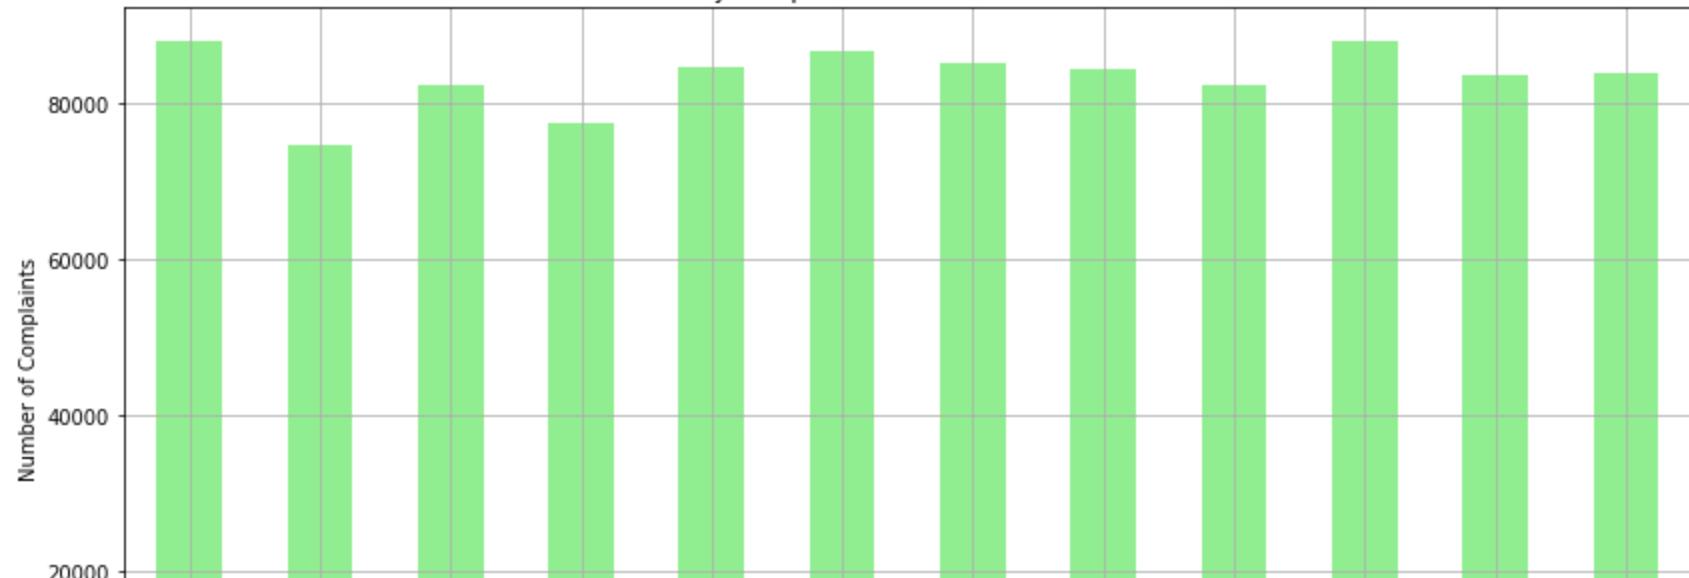
# Monthly Volumes
monthly_volume.plot(kind='bar', ax=ax[1], color='lightgreen')
ax[1].set_title('Monthly Complaint Volumes (Across All Years)')
ax[1].set_xlabel('Month')
ax[1].set_ylabel('Number of Complaints')
ax[1].grid()

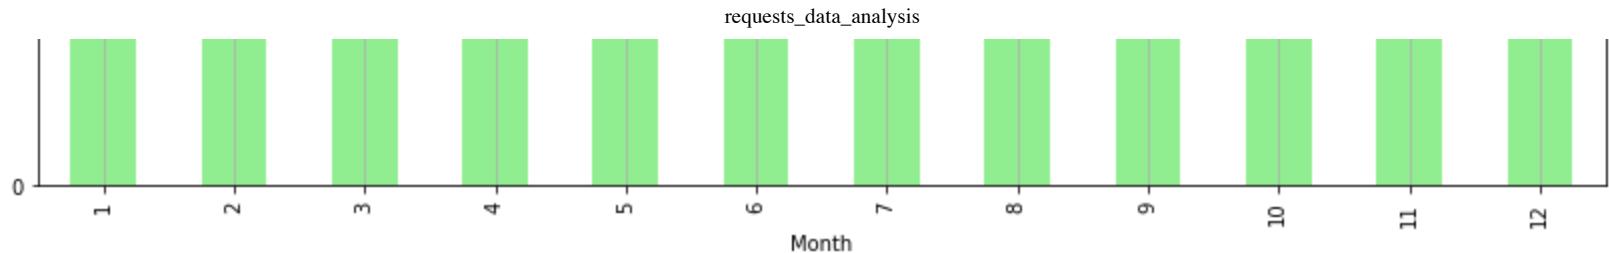
plt.tight_layout()
plt.show()
```

Yearly Complaint Volumes



Monthly Complaint Volumes (Across All Years)





In [5]:

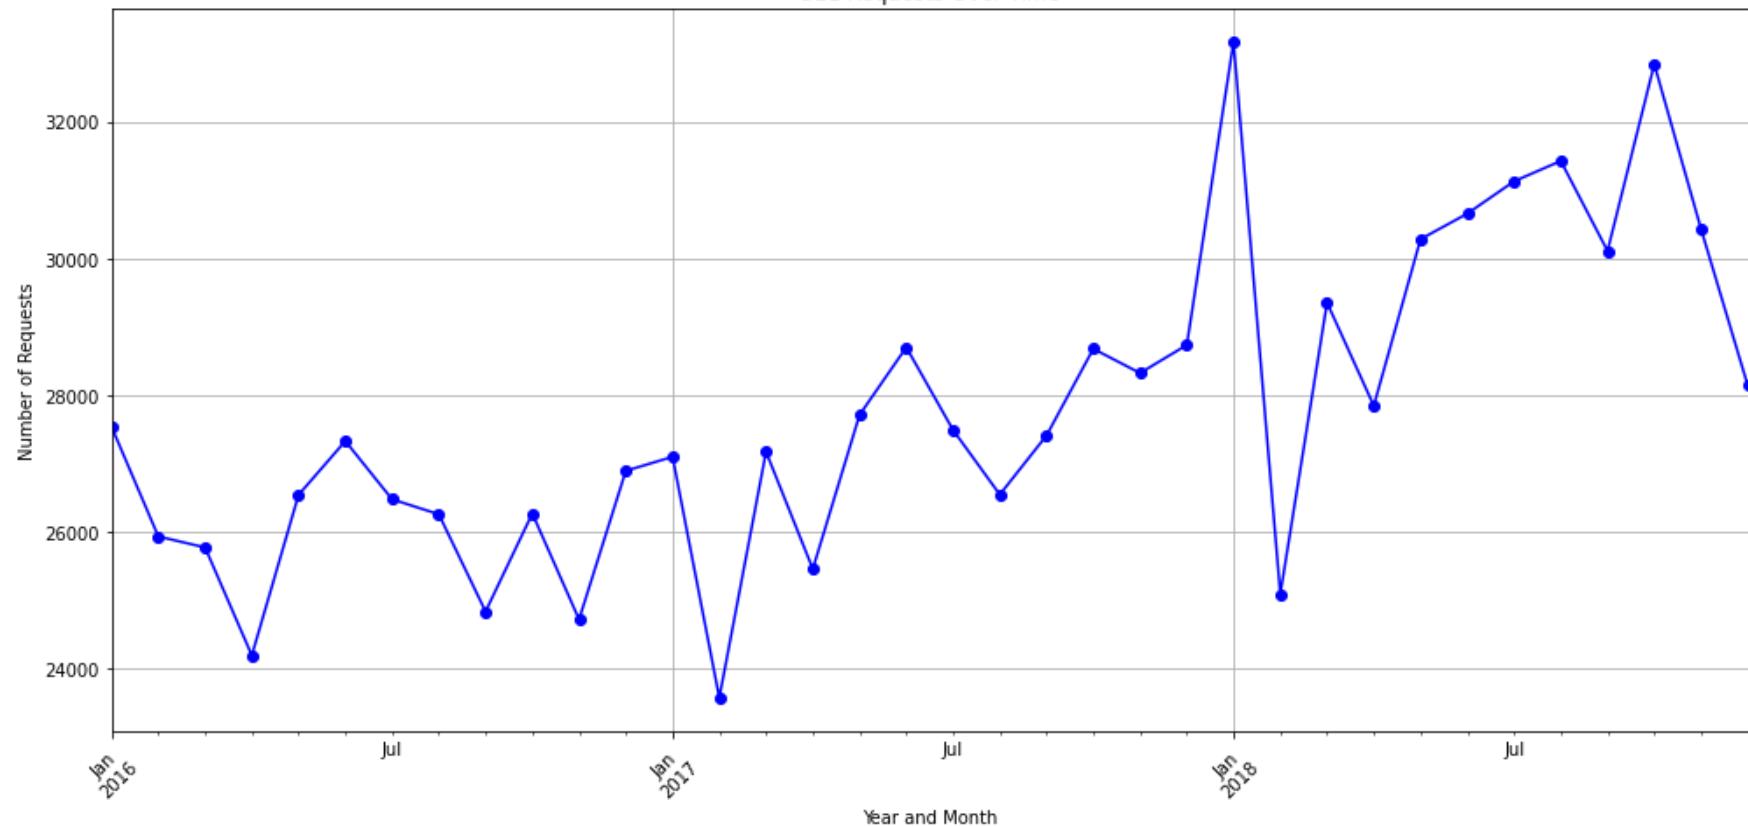
```
# Extract year and month for grouping purposes
requests['Year-Month'] = requests['Created Date'].dt.to_period('M')

# Group by the new 'Year-Month' column to see the number of requests over time
requests_by_month = requests.groupby('Year-Month').size()

# Plotting
plt.figure(figsize=(14, 7))
requests_by_month.plot(kind='line', color='blue', marker='o')
plt.title('311 Requests Over Time')
plt.xlabel('Year and Month')
plt.ylabel('Number of Requests')
plt.xticks(rotation=45)
plt.grid(True)

plt.tight_layout()
plt.show()
```

311 Requests Over Time



The visualizations provide insights into the inbound call patterns based on the 311 records:

- 1. Yearly Complaint Volumes:** The bar chart shows the number of complaints for each year represented in the dataset. It looks like the number of complaints have increased over the 3 years.
- 2. Monthly Complaint Volumes (Across All Years):** The second bar chart aggregates complaint volumes by month, combining data from all years. Not really seeing any major trends here - this is somewhat surprising since I expected seasonal trends.

We can also do this weekly.

In [6]:

```
# Extract the week of the year from the 'Request Date'
requests['Request Week'] = requests['Created Date'].dt.isocalendar().week

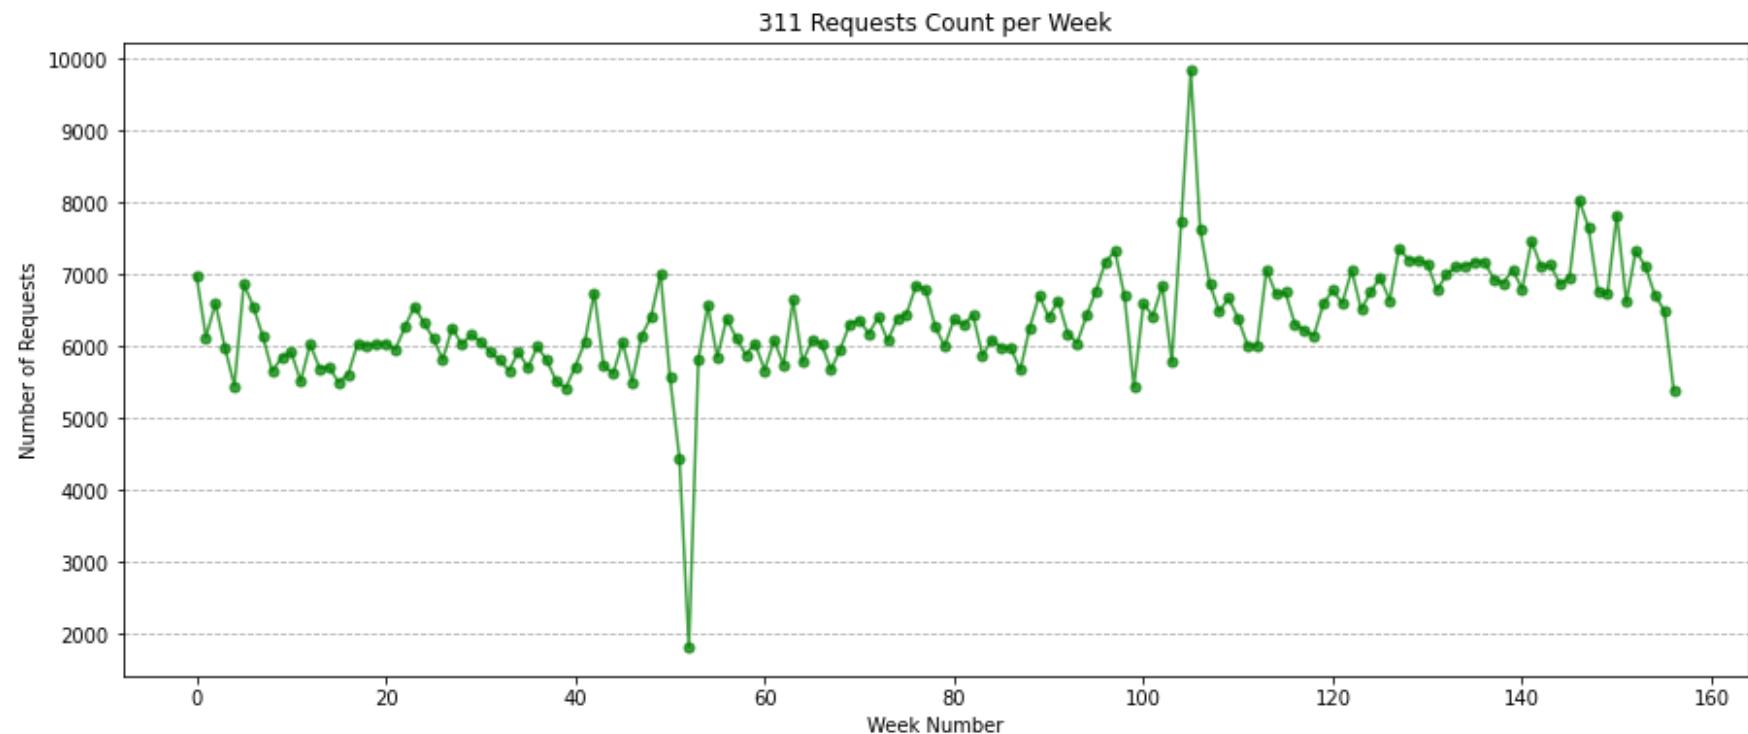
# Group by year as well to ensure accuracy across years
requests['Request Year'] = requests['Created Date'].dt.year
```

```
# Aggregate the data by year and week
weekly_requests = requests.groupby(['Request Year', 'Request Week']).size()

# Reset index to have a flat structure for easier plotting
weekly_requests_reset = weekly_requests.reset_index(name='Count')

# Plotting the weekly request counts
plt.figure(figsize=(15, 6))
plt.plot(weekly_requests_reset.index, weekly_requests_reset['Count'], marker='o', linestyle='--', markersize=5,
plt.title('311 Requests Count per Week')
plt.xlabel('Week Number')
plt.ylabel('Number of Requests')
plt.grid(axis='y', linestyle='--')

plt.show()
```



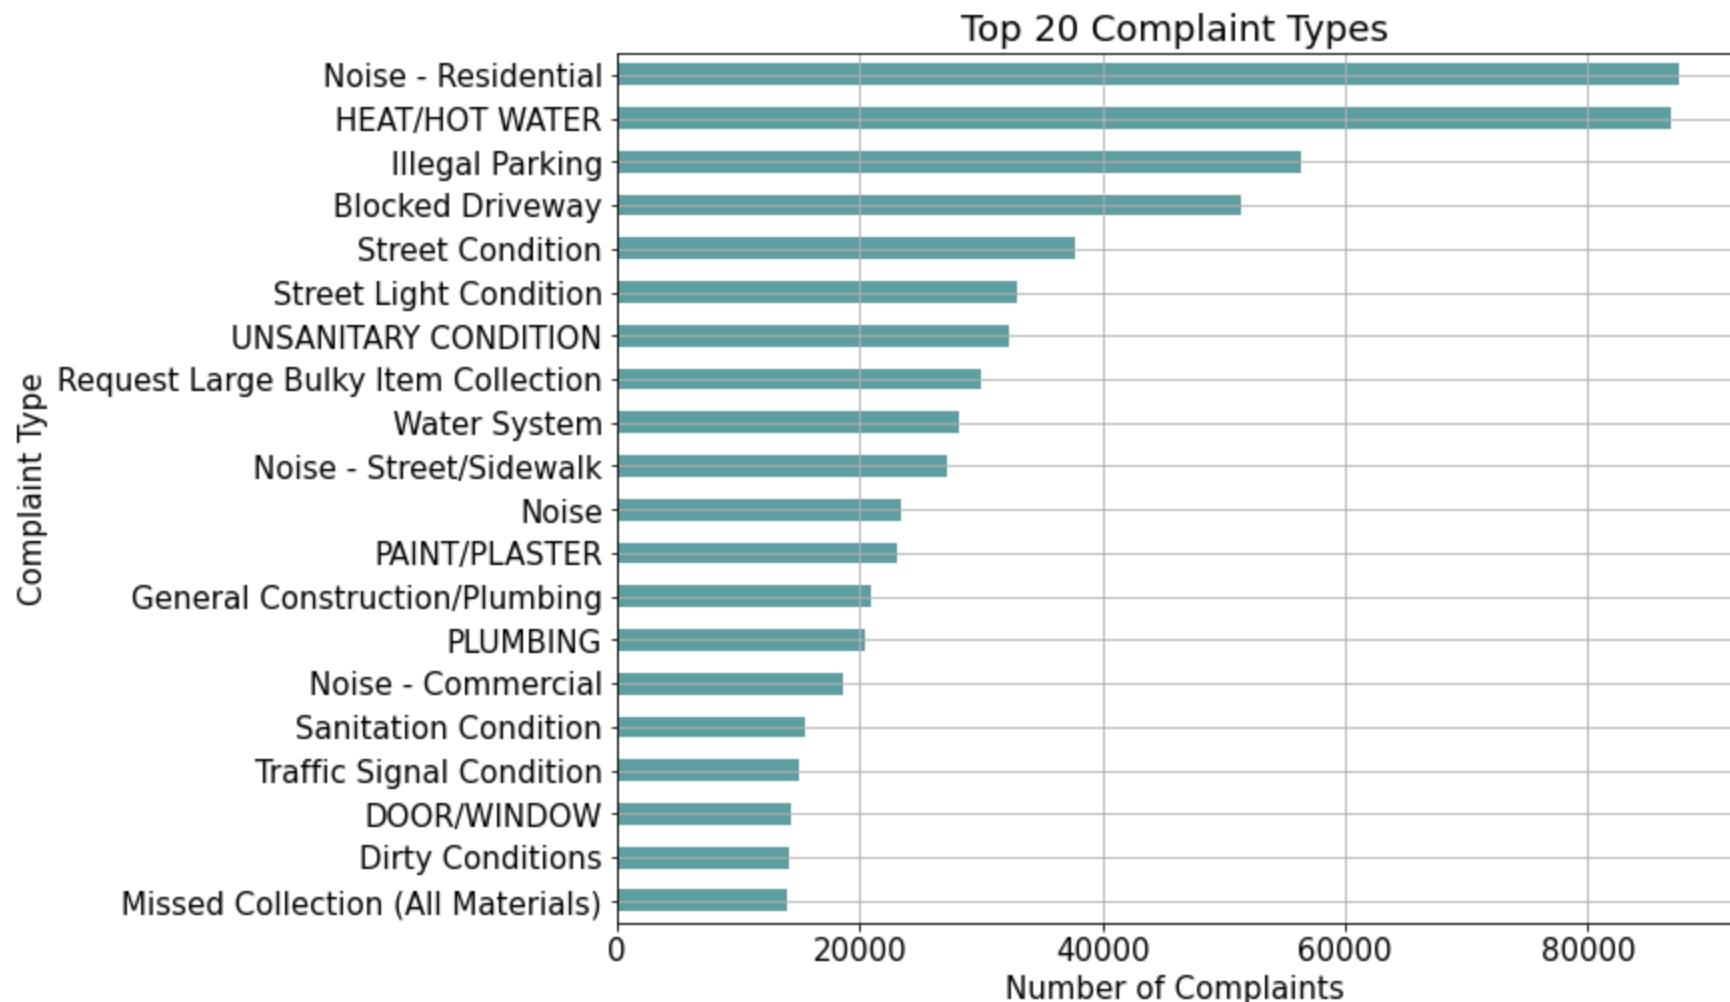
Complaint Type Distribution

We'll next look at the distribution of complaint types. This will help us understand the most common issues reported by residents. After that, we can explore the resolution times to gauge the efficiency of the agencies' responses.

```
In [7]: complaint_type_distribution = requests['Complaint Type'].value_counts().head(20)

# Plotting
font = {'size' : 15}
plt.rc('font', **font)

plt.figure(figsize=(10, 8))
complaint_type_distribution.plot(kind='barh', color='cadetblue')
plt.title('Top 20 Complaint Types')
plt.xlabel('Number of Complaints')
plt.ylabel('Complaint Type')
plt.gca().invert_yaxis() # To display the highest number at the top
plt.grid()
plt.show()
```



The bar chart displays the top 20 complaint types based on their frequencies. This visualization helps us identify the most common issues that residents' report. As we can see, **Noise - Residential** and **Heat/Hot Water** are the most frequent types of requests.

Agency Name Distribution

We'll next look at the distribution of frequency of complaints by agency.

In [8]:

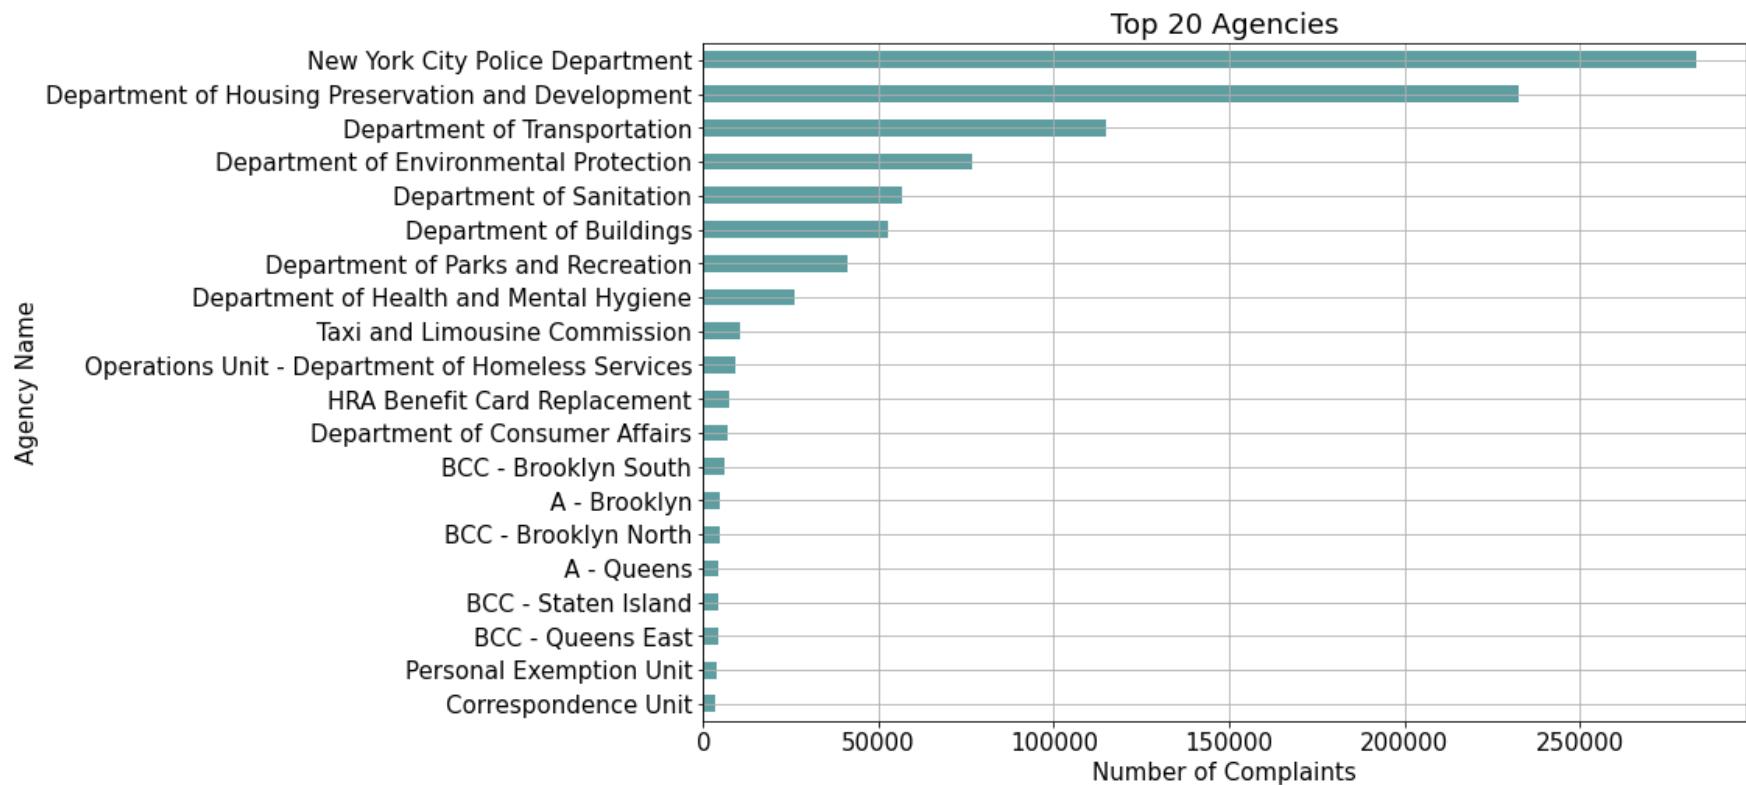
```
agency_name_distribution = requests['Agency Name'].value_counts().head(20)

plt.figure(figsize=(12, 8))
agency_name_distribution.plot(kind='barh', color='cadetblue')
```

```

plt.title('Top 20 Agencies')
plt.xlabel('Number of Complaints')
plt.ylabel('Agency Name')
plt.gca().invert_yaxis() # To display the highest number at the top
plt.grid()
plt.show()

```



We see that NYCPD and Department of Housing Preservation and Development receive the highest number of complaints.

Which Borough has the highest requests?

In [9]:

```

import matplotlib.pyplot as plt

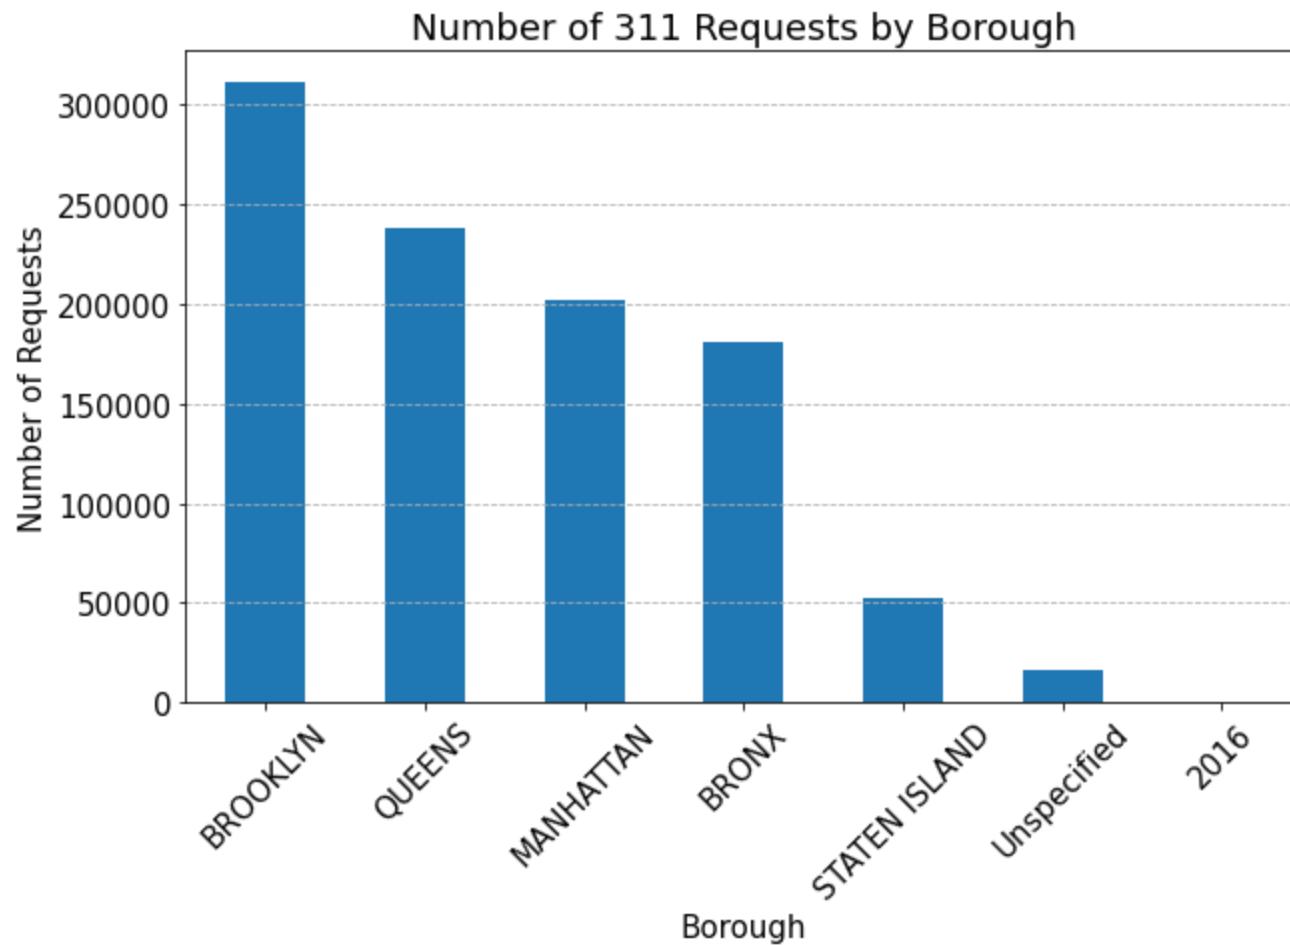
# Count the number of complaints/requests from different boroughs
borough_counts = requests['Borough'].value_counts()

# Plotting the distribution of complaints/requests across different boroughs
plt.figure(figsize=(10, 6))
borough_counts.plot(kind='bar')

```

```
plt.title('Number of 311 Requests by Borough')
plt.xlabel('Borough')
plt.ylabel('Number of Requests')
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--')

plt.show()
```



Given that the different boroughs have different populations, we can get the per capita complaint frequency by normalizing the above counts by the population. The population of different Boroughs from a recent census is as follows -

Brooklyn: 2.6 million Queens: 2.3 million Manhattan: 1.6 million The Bronx: 1.4 million Staten Island: 0.5 million

Next, we'll calculate the per capita complaints/requests for each borough by dividing the total number of complaints/requests from each borough by its population.

In [10]:

```
# Correcting the calculation for per capita requests

# Population data for each borough
populations = {
    'BROOKLYN': 2.6e6,
    'QUEENS': 2.3e6,
    'MANHATTAN': 1.6e6,
    'BRONX': 1.4e6,
    'STATEN ISLAND': 0.5e6
}

# First, ensure we correctly aggregate the number of requests per borough according to the updated population !
borough_counts_corrected = requests['Borough'].value_counts()

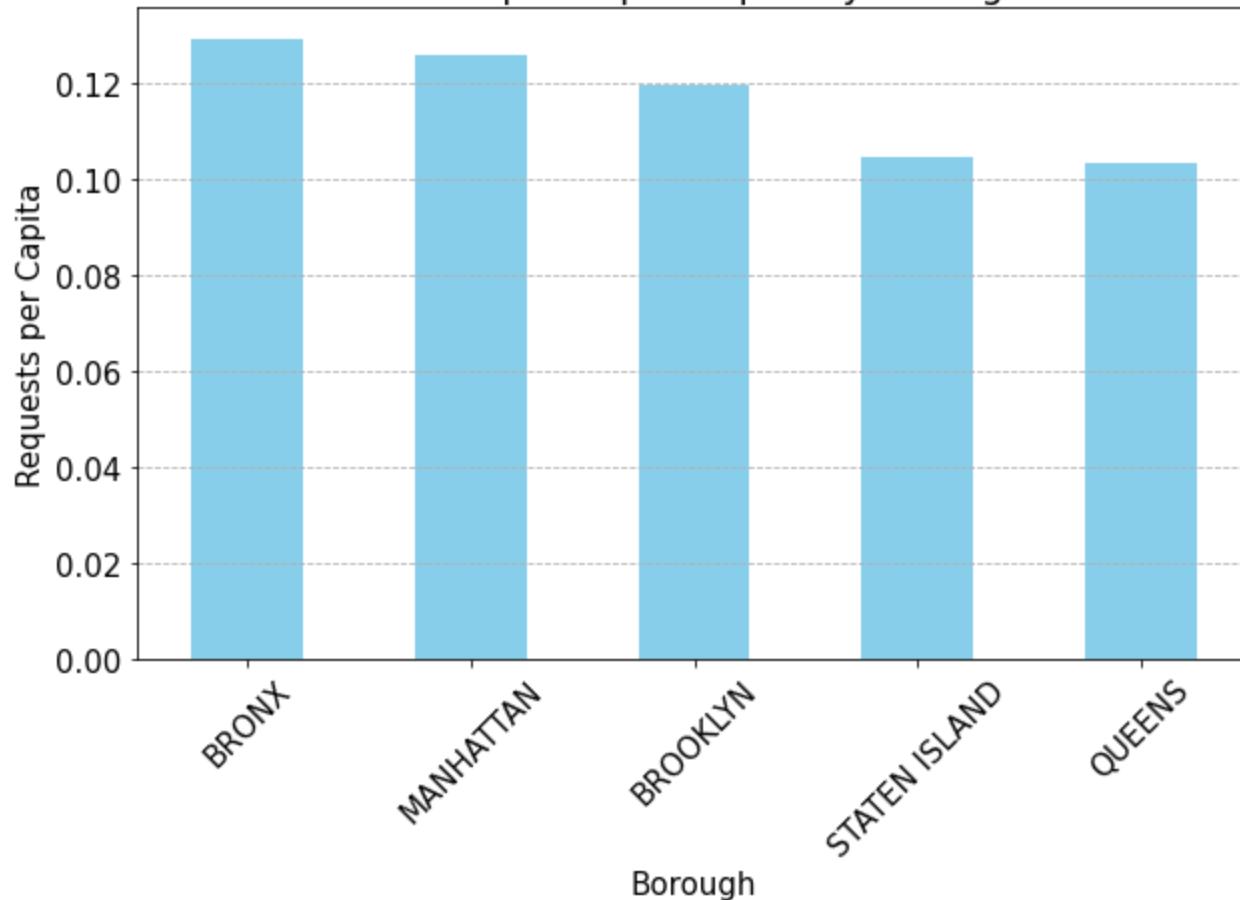
# Calculate the per capita complaints/requests for each borough using the corrected aggregation
per_capita_requests_corrected = {borough: borough_counts_corrected[borough] / populations[borough] for borough in populations}

per_capita_requests_corrected_df = pd.Series(per_capita_requests_corrected).sort_values(ascending=False)

# Plotting the per capita complaints/requests by borough
plt.figure(figsize=(10, 6))
per_capita_requests_corrected_df.plot(kind='bar', color='skyblue')
plt.title('311 Requests per Capita by Borough')
plt.xlabel('Borough')
plt.ylabel('Requests per Capita')
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--')

plt.show()
```

311 Requests per Capita by Borough



Geographical locations of requests

Since we have the latitude, longitude of the reported incident, we can plot the locations and density on a map of New York (mostly because it's cool!).

In [11]:

```
requests_lat_long = requests.copy()

requests_lat_long['Latitude'] = np.round(requests_lat_long['Latitude'],2)
requests_lat_long['Longitude'] = np.round(requests_lat_long['Longitude'],2)

requests_lat_long_agg = requests_lat_long.groupby(['Latitude','Longitude']).size().reset_index(name='counts')
requests_lat_long_agg['Density'] = requests_lat_long_agg['counts']/requests_lat_long_agg['counts'].sum()
```

```
color_scale = [(0, 'orange'), (1, 'red')]

fig = px.density_mapbox(requests_lat_long_agg,
                        lat="Latitude",
                        lon="Longitude",
                        z="Density",
                        color_continuous_scale='viridis',
                        zoom=9,
                        radius=26,
                        height=800,
                        width=900)

fig.update_layout(mapbox_style="open-street-map")
fig.show()
```





Next, let's explore the resolution times for complaints to understand how quickly agencies respond to and resolve these issues. We'll calculate the resolution time as the difference between the 'Closed Date' and 'Created Date' and then analyze the average resolution time by complaint type. Since we have a large number of different types of complaints, we will look at the 20 complaints that have the fastest and slowest resolution times.

In [12]:

```
requests['Created Date'] = pd.to_datetime(requests['Created Date'], errors='coerce')
requests['Closed Date'] = pd.to_datetime(requests['Closed Date'], errors='coerce')

# Calculate resolution time in days
requests['Resolution Time'] = (requests['Closed Date'] - requests['Created Date']).dt.total_seconds() / (60 * 60)

# Group by 'Complaint Type' and calculate average resolution time
average_resolution_time = requests.groupby('Complaint Type')['Resolution Time'].mean().reset_index()

average_resolution_time.sort_values(by='Resolution Time', ascending=False)
```

Out [12]:

	Complaint Type	Resolution Time
233	Unsanitary Condition	856.811407
105	Found Property	600.863376
108	General	527.235428

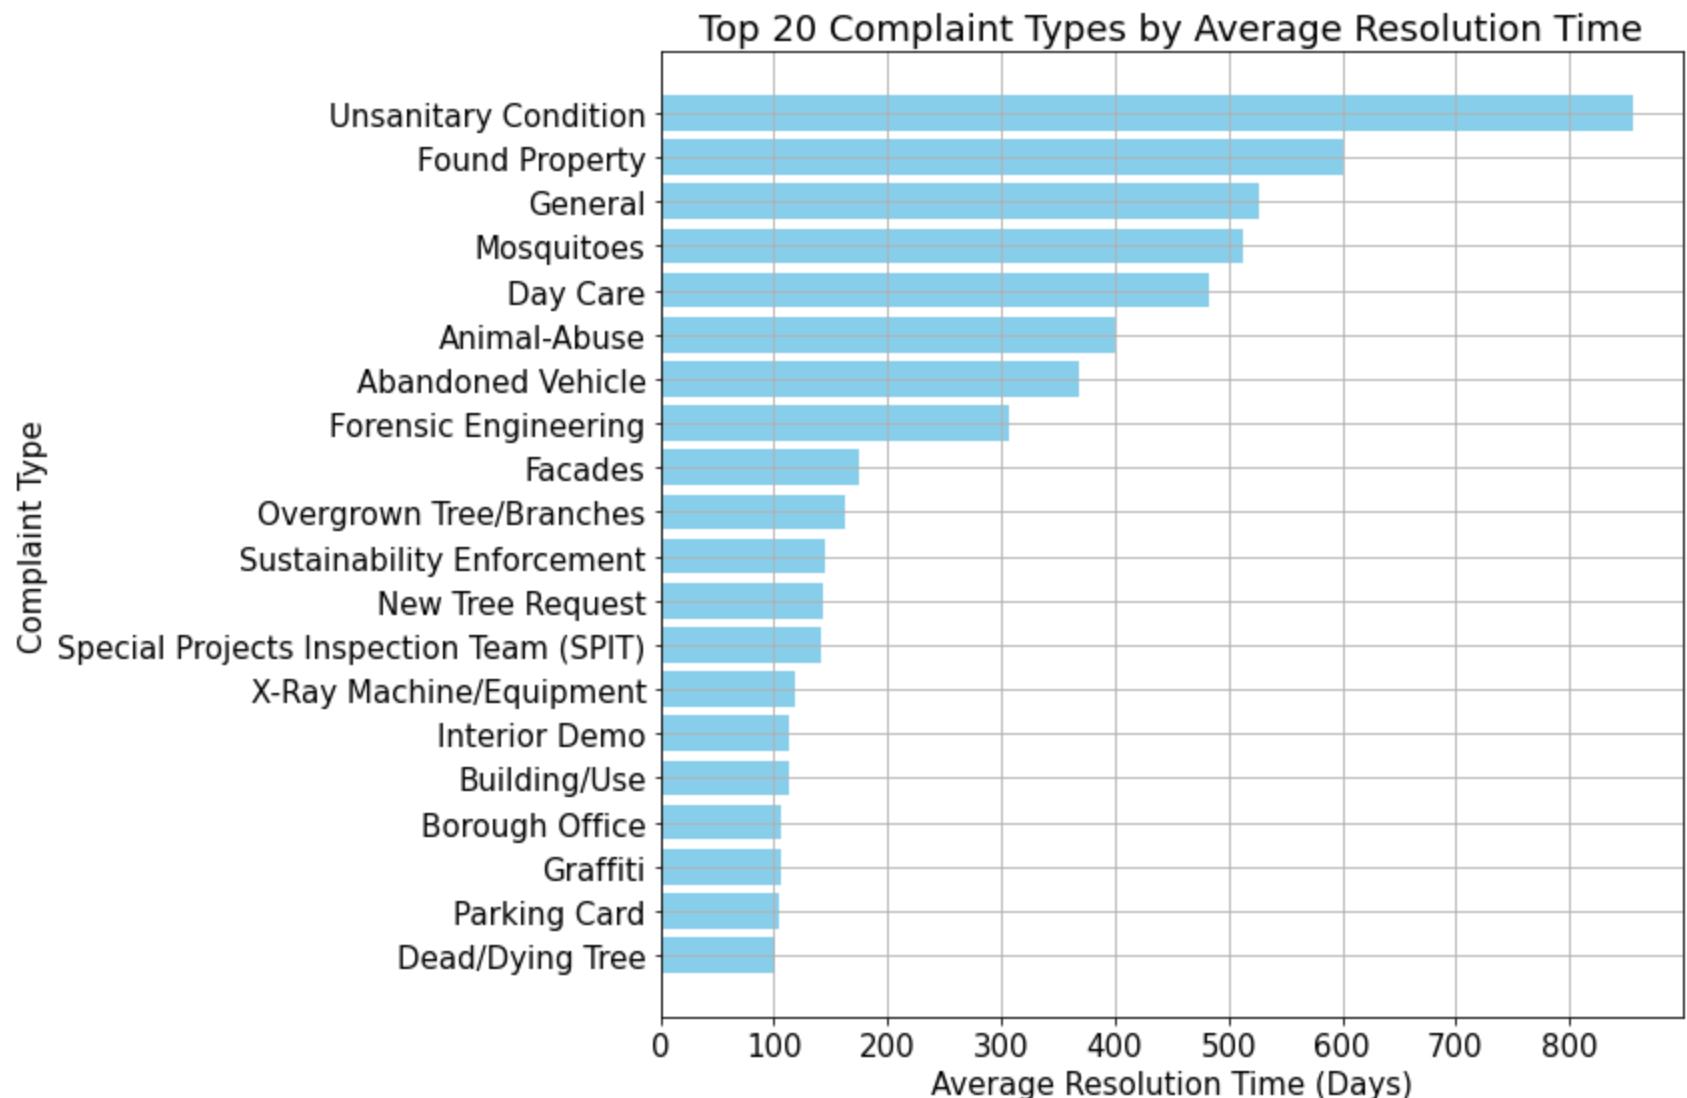
	Complaint Type	Resolution Time
147	Mosquitoes	512.553917
68	Day Care	483.078656
...
225	Trans Fat	NaN
230	Unlicensed Dog	NaN
231	Unsanitary Animal Facility	NaN
236	VACANT APARTMENT	NaN
244	Window Guard	NaN

247 rows × 2 columns

In [13]:

```
# Since the list is long, let's take the top 20 complaint types for a clearer visualization
top_20_average_resolution_time = average_resolution_time.nlargest(20, 'Resolution Time')

# Plotting
plt.figure(figsize=(12, 8))
plt.barh(top_20_average_resolution_time['Complaint Type'], top_20_average_resolution_time['Resolution Time'],
        plt.xlabel('Average Resolution Time (Days)')
        plt.ylabel('Complaint Type')
        plt.title('Top 20 Complaint Types by Average Resolution Time')
        plt.gca().invert_yaxis() # To display the longest resolution time at the top
        plt.tight_layout()
        plt.grid()
        plt.show()
```

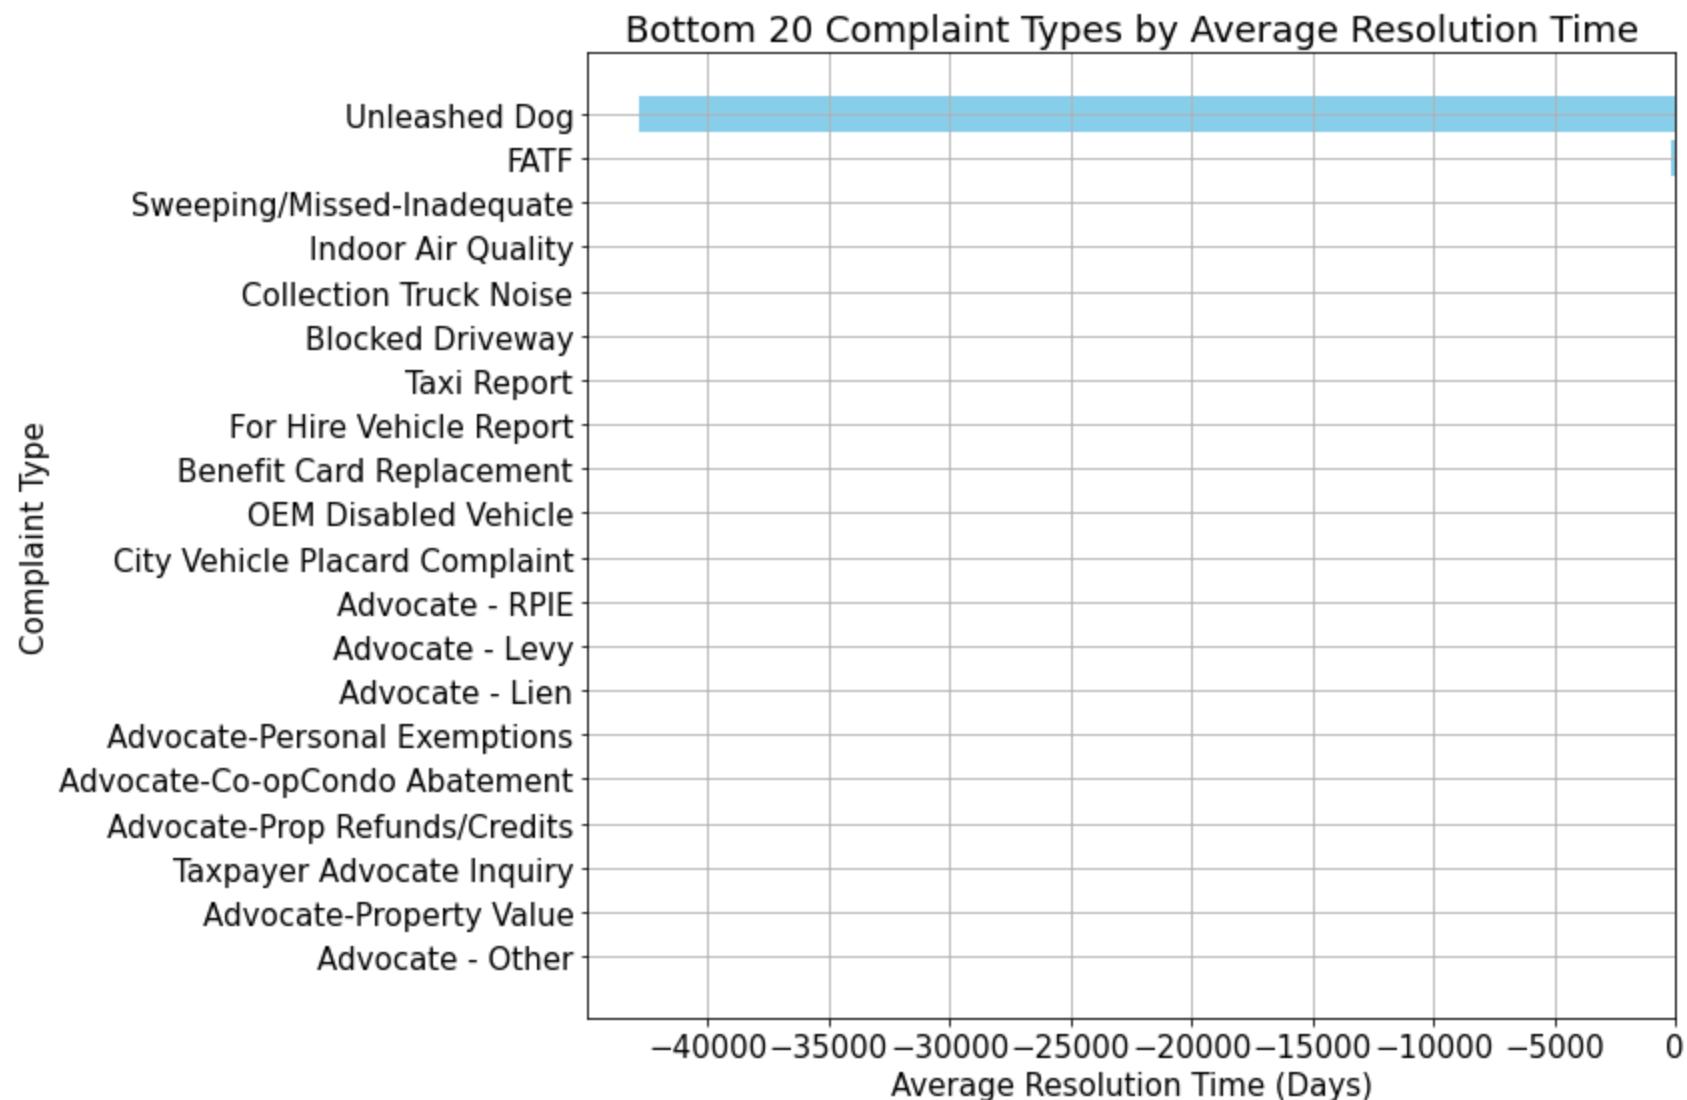


In [14]:

```
# Since the list is long, let's take the top 20 complaint types for a clearer visualization
bottom_20_average_resolution_time = average_resolution_time.dropna().nsmallest(20, 'Resolution Time')

# Plotting
plt.figure(figsize=(12, 8))
plt.barh(bottom_20_average_resolution_time['Complaint Type'], bottom_20_average_resolution_time['Resolution Time'])
plt.xlabel('Average Resolution Time (Days)')
plt.ylabel('Complaint Type')
plt.title('Bottom 20 Complaint Types by Average Resolution Time')
```

```
plt.gca().invert_yaxis() # To display the longest resolution time at the top  
plt.tight_layout()  
plt.grid()  
plt.show()
```



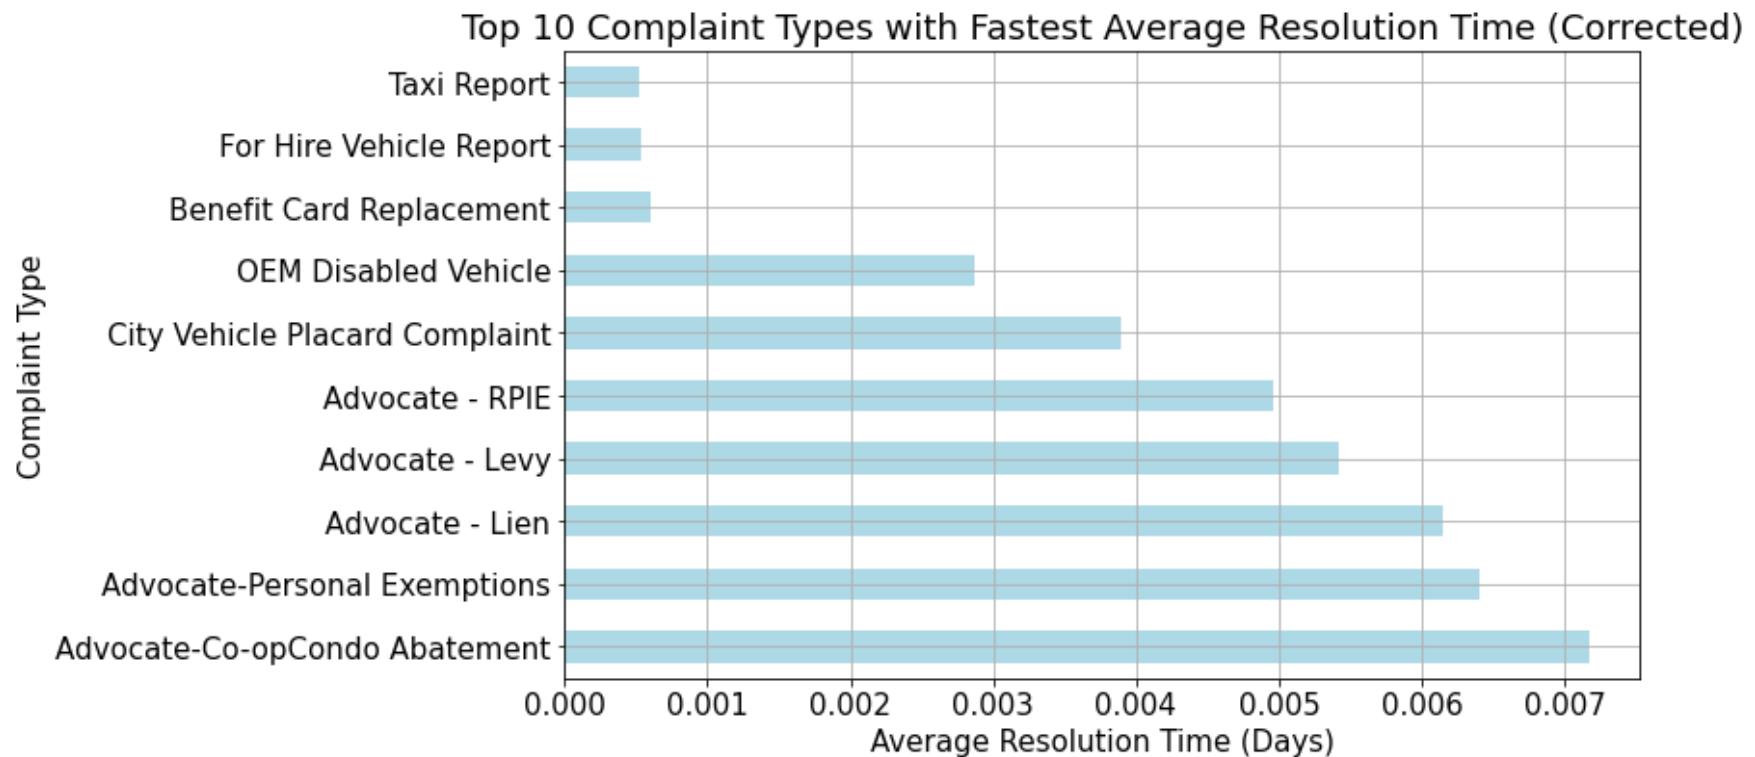
So we're seeing some negative resolution times, which is an error in the data. To address this issue of negative resolution times, which likely represent errors in the data (such as closed dates recorded before the created dates), we will first filter out these erroneous records. After correcting for negative resolution times, we'll re-analyze the average resolution times by complaint type, ensuring more accurate insights.

In [15]:

```
# Filter out records with negative resolution times
data_corrected = requests[requests['Resolution Time'] >= 0]

# Re-calculate the average resolution time by complaint type, focusing on corrected data
average_resolution_time_corrected = data_corrected.groupby('Complaint Type')['Resolution Time'].mean().sort_values(ascending=True)

# Plotting the corrected average resolution times
plt.figure(figsize=(10, 6))
average_resolution_time_corrected.plot(kind='barh', color='lightblue')
plt.title('Top 10 Complaint Types with Fastest Average Resolution Time (Corrected)')
plt.xlabel('Average Resolution Time (Days)')
plt.ylabel('Complaint Type')
plt.gca().invert_yaxis() # To display the lowest number at the top
plt.grid()
plt.show()
```



Okay that looks much better. Although we are seeing really quick resolution times (a few seconds), this is likely because some of these were marked as resolved on the initial phone call itself.

Weather Data

Now we'll jump into the weather data and start exploring this dataset.

In [16]:

```
weather_data = pd.read_csv('weather_NY_2010_2018Nov.csv')
weather_data['Precipitation'] = weather_data['Percipitation']

weather_data['dt'] = pd.to_datetime(weather_data.Year.astype(str) + '/' + weather_data.Month.astype(str) + '/' + weather_data.Day.astype(str))
weather_data
```

Out[16]:

	USAF	WBAN	StationName	State	Latitude	Longitude	MeanTemp	MinTemp	MaxTemp	DewPoint	...	MaxSustainedWind	...
0	726228	94740	ADIRONDACK REGIONAL ARPT	NY	44.385	-74.207	27.6	24.8	30.9	25.0	...	6.0	
1	726228	94740	ADIRONDACK REGIONAL ARPT	NY	44.385	-74.207	-3.2	-20.9	17.1	-9.6	...	9.9	
2	726228	94740	ADIRONDACK REGIONAL ARPT	NY	44.385	-74.207	20.9	17.1	24.1	15.1	...	12.0	
3	726228	94740	ADIRONDACK REGIONAL ARPT	NY	44.385	-74.207	13.8	5.0	19.9	8.5	...	8.0	
4	726228	94740	ADIRONDACK REGIONAL ARPT	NY	44.385	-74.207	6.3	-8.0	19.0	1.9	...	5.1	
...
160770	743700	14715	WHEELER-SACK AAF AIRPORT	NY	44.050	-75.733	78.4	68.5	90.9	71.6	...	13.0	
160771	743700	14715	WHEELER-SACK AAF AIRPORT	NY	44.050	-75.733	75.3	66.2	82.0	71.4	...	9.9	
160772	743700	14715	WHEELER-SACK AAF AIRPORT	NY	44.050	-75.733	63.0	51.4	73.2	54.0	...	8.0	

USAF	WBAN	StationName	State	Latitude	Longitude	MeanTemp	MinTemp	MaxTemp	DewPoint	...	MaxSustainedWind	Gust	Rain	SnowIce	SnowDepth	Year	Month	Day
160773	743700	14715	WHEELER-SACK AAF AIRPORT	NY	44.050	-75.733	56.4	49.5	73.2	45.7	...	14.0						
160774	743700	14715	WHEELER-SACK AAF AIRPORT	NY	44.050	-75.733	51.3	40.6	62.6	43.4	...	8.9						

160775 rows × 22 columns

The weather dataset contains 160,775 entries, with 20 columns detailing weather observations from New York between 2010 and November 2018. Here's a brief overview of some key columns:

- USAF, WBAN: Identifiers for the weather stations.
- StationName: The name of the weather station.
- State: The state of the weather station (all NY in this dataset).
- Latitude, Longitude: Geographic coordinates of the weather station.
- MeanTemp: The mean temperature for the day.
- MinTemp, MaxTemp: The minimum and maximum temperatures for the day.
- DewPoint: The dew point temperature.
- Precipitation: The amount of precipitation.
- WindSpeed, MaxSustainedWind: Wind speed and maximum sustained wind speed.
- Gust: Wind gust speed.
- Rain, SnowIce: Indicators for the presence of rain or snow/ice.
- SnowDepth: The depth of snow.
- Year, Month, Day: The date of the observation.

For basic exploratory data analysis with plots, we could look into the following:

- Temperature Trends: Analyze how temperatures (mean, min, max) have changed over the years.
- Precipitation Patterns: Explore the distribution and trends of precipitation over time.
- Wind Patterns: Examine wind speed and gusts over the years.

In [17]:

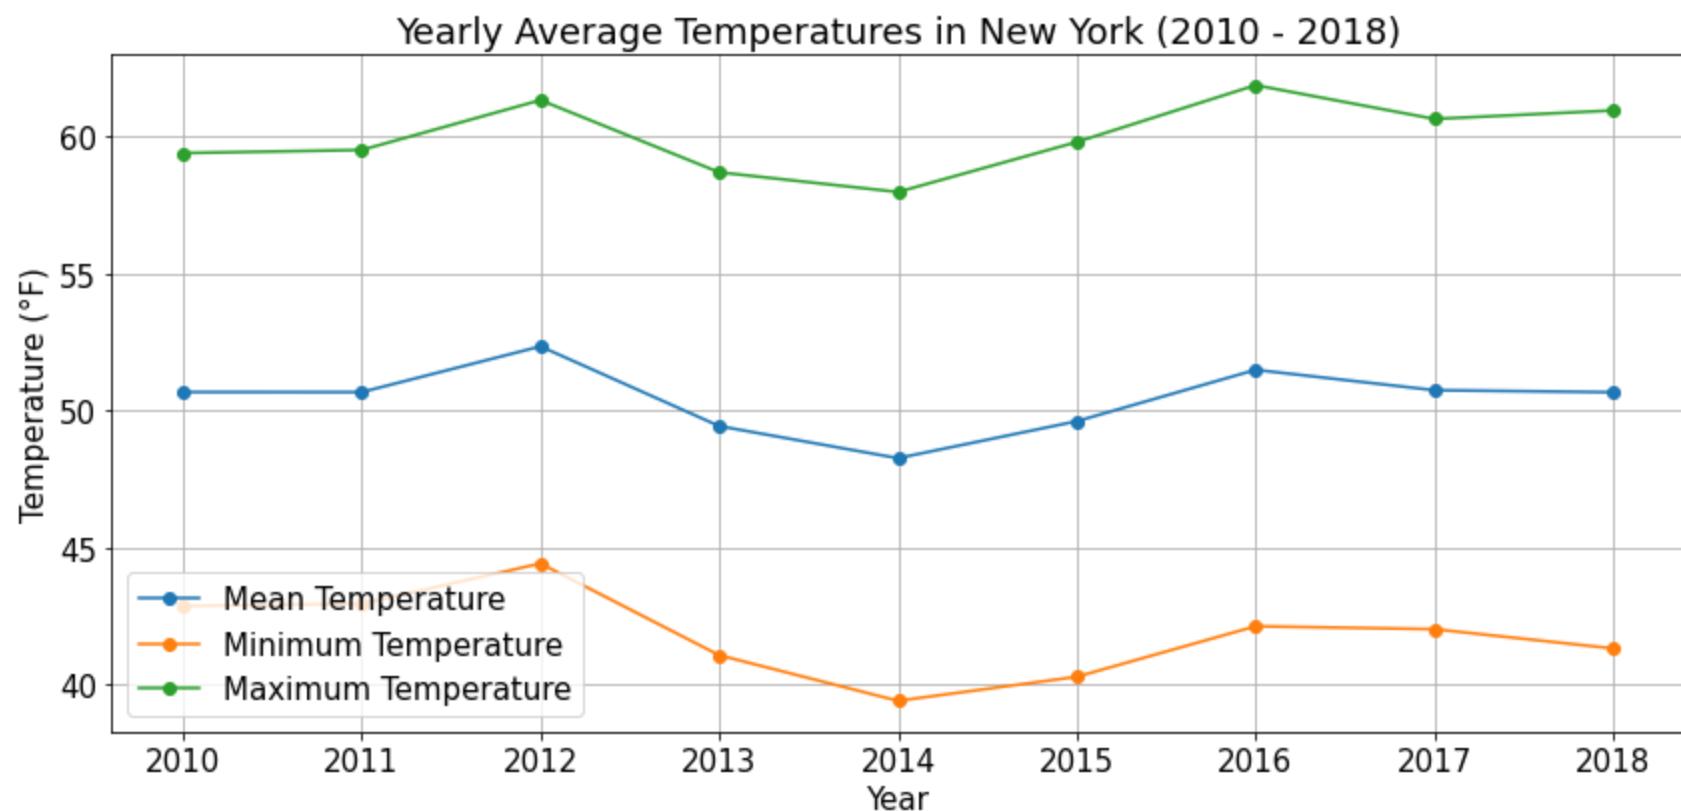
```
# Calculate yearly average temperatures
yearly_temps = weather_data.groupby('Year').agg({'MeanTemp': 'mean', 'MinTemp': 'mean', 'MaxTemp': 'mean'})
```

```
# Plotting
plt.figure(figsize=(12, 6))

plt.plot(yearly_temps.index, yearly_temps['MeanTemp'], label='Mean Temperature', marker='o')
plt.plot(yearly_temps.index, yearly_temps['MinTemp'], label='Minimum Temperature', marker='o')
plt.plot(yearly_temps.index, yearly_temps['MaxTemp'], label='Maximum Temperature', marker='o')

plt.title('Yearly Average Temperatures in New York (2010 - 2018)')
plt.xlabel('Year')
plt.ylabel('Temperature (°F)')
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()
```



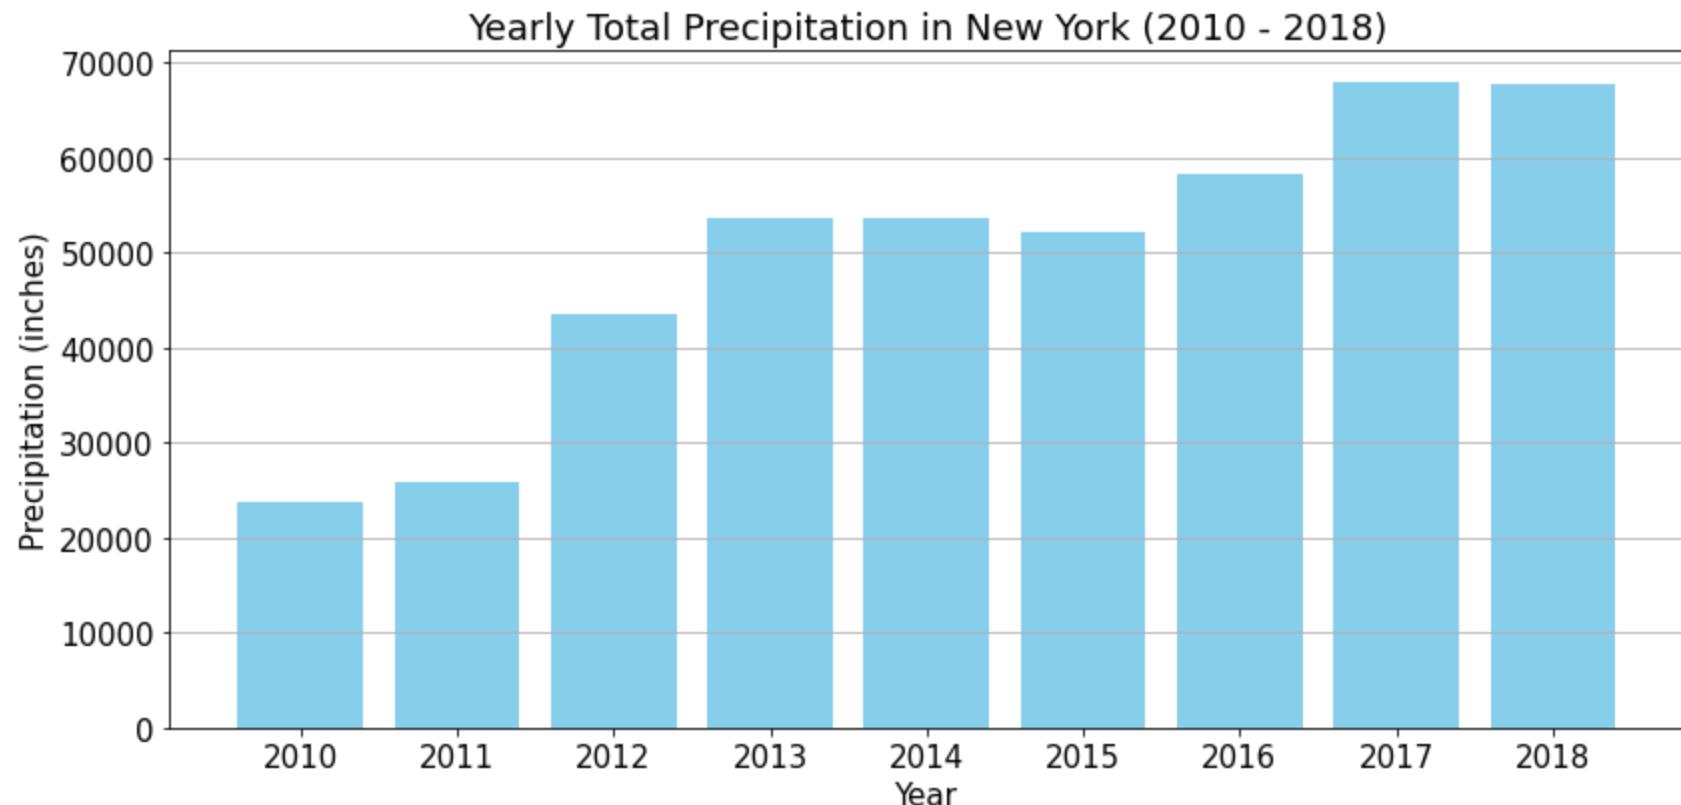
The plot illustrates the yearly average temperatures (mean, minimum, and maximum) in New York from 2010 through 2018. Each line represents the trend for mean, minimum, and maximum temperatures, providing insight into how the climate has varied over these

years.

In [18]:

```
# Calculate yearly total precipitation
yearly_precipitation = weather_data.groupby('Year')['Percipitation'].sum()

# Plotting the yearly total precipitation
plt.figure(figsize=(12, 6))
plt.bar(yearly_precipitation.index, yearly_precipitation.values, color='skyblue')
plt.title('Yearly Total Precipitation in New York (2010 - 2018)')
plt.xlabel('Year')
plt.ylabel('Precipitation (inches)')
plt.xticks(yearly_precipitation.index)
plt.grid(axis='y')
plt.tight_layout()
plt.show()
```



Looks like NY has received increased amounts of precipitation over the years.

The bar chart displays the yearly total precipitation in New York from 2010 through 2018. This visualization helps identify the years with higher or lower amounts of precipitation, indicating variability in rainfall and snowfall over the years.

In [19]:

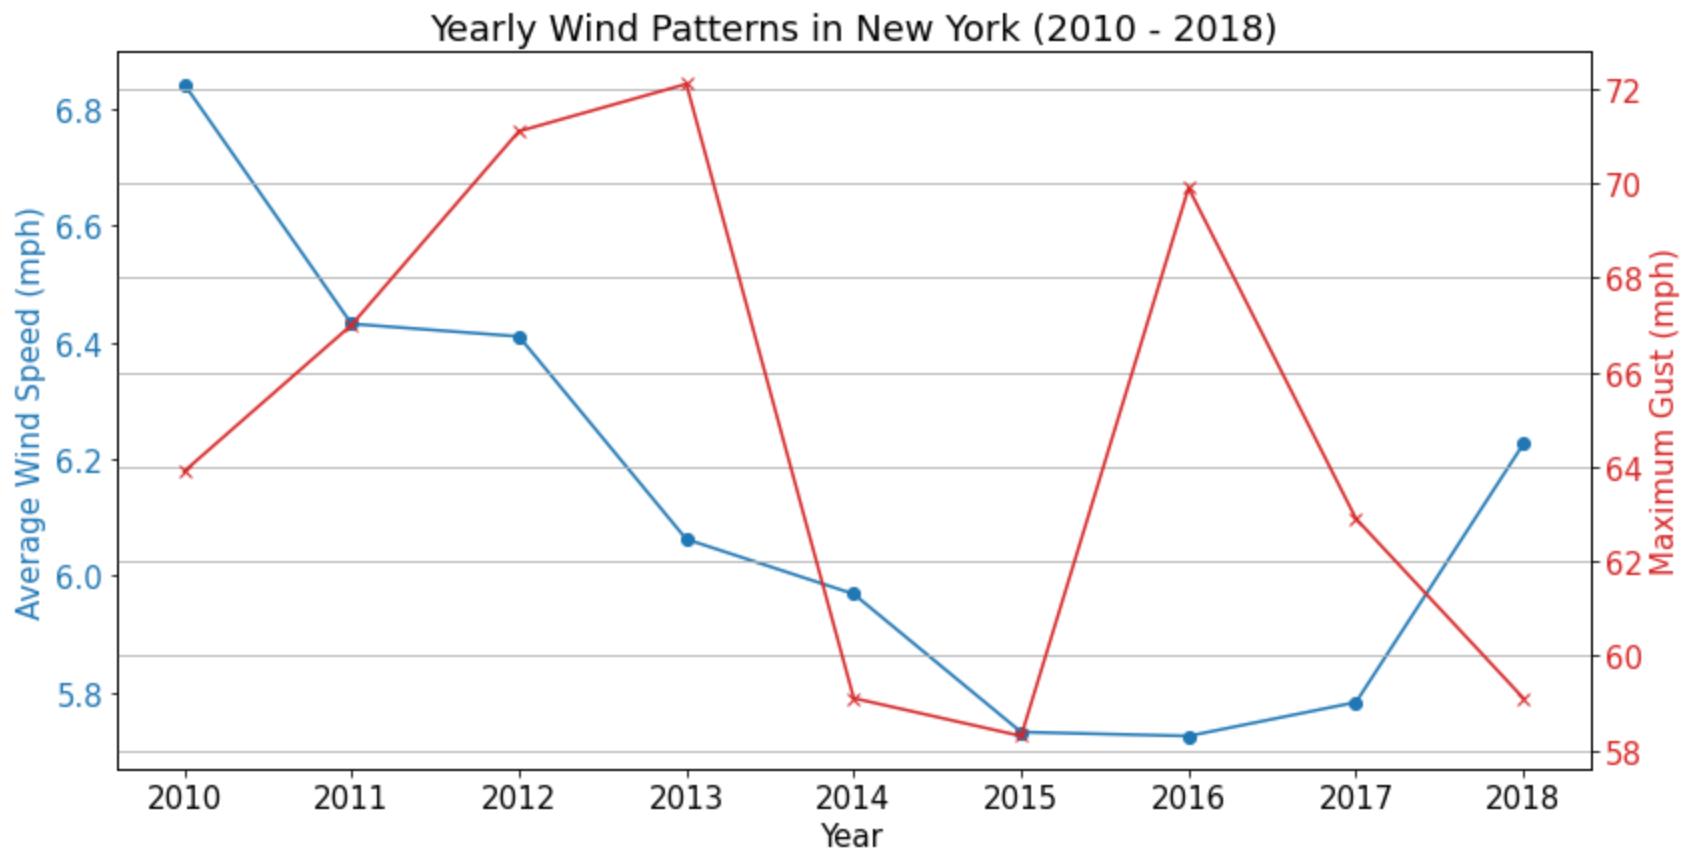
```
# Calculate yearly average wind speed and maximum gust
yearly_wind = weather_data.groupby('Year').agg({'WindSpeed': 'mean', 'Gust': 'max'})

# Plotting
fig, ax1 = plt.subplots(figsize=(12, 6))

color = 'tab:blue'
ax1.set_xlabel('Year')
ax1.set_ylabel('Average Wind Speed (mph)', color=color)
ax1.plot(yearly_wind.index, yearly_wind['WindSpeed'], label='Average Wind Speed', color=color, marker='o')
ax1.tick_params(axis='y', labelcolor=color)

ax2 = ax1.twinx() # instantiate a second axes that shares the same x-axis
color = 'tab:red'
ax2.set_ylabel('Maximum Gust (mph)', color=color) # we already handled the x-label with ax1
ax2.plot(yearly_wind.index, yearly_wind['Gust'], label='Maximum Gust', color=color, marker='x')
ax2.tick_params(axis='y', labelcolor=color)

fig.tight_layout() # otherwise the right y-label is slightly clipped
plt.title('Yearly Wind Patterns in New York (2010 – 2018)')
plt.grid(True)
plt.show()
```



The chart above illustrates the yearly wind patterns in New York from 2010 through 2018, showing both the average wind speed (in mph) and the maximum gusts recorded each year. The blue line represents the average wind speed, while the red line indicates the maximum gust speeds.

This visualization provides insight into the general wind conditions and highlights years with particularly strong gusts, reflecting the variability of wind intensity over the observed period.

Given the analyses conducted so far on the weather dataset, including temperature trends, precipitation patterns, and wind patterns, another insightful exploration could involve examining the relationship between weather conditions and specific weather events, such as days with significant precipitation or high wind speeds, and their occurrence over the years.

For this next step, let's focus on:

1. Days with Significant Precipitation: Identify and visualize the number of days per year with precipitation exceeding a certain threshold, indicating heavy rainfall or significant snowfall events.\
2. Days with High Wind Speeds: Analyze and plot the number of days per year with wind speeds exceeding a threshold, highlighting windy conditions or storms.

Days with Significant Precipitation

Let's define "significant precipitation" as days with precipitation amounts exceeding 0.5 inches, which often indicates heavy rainfall or significant snowfall events. We'll count these days per year to see any trends in their frequency over time.

In [20]:

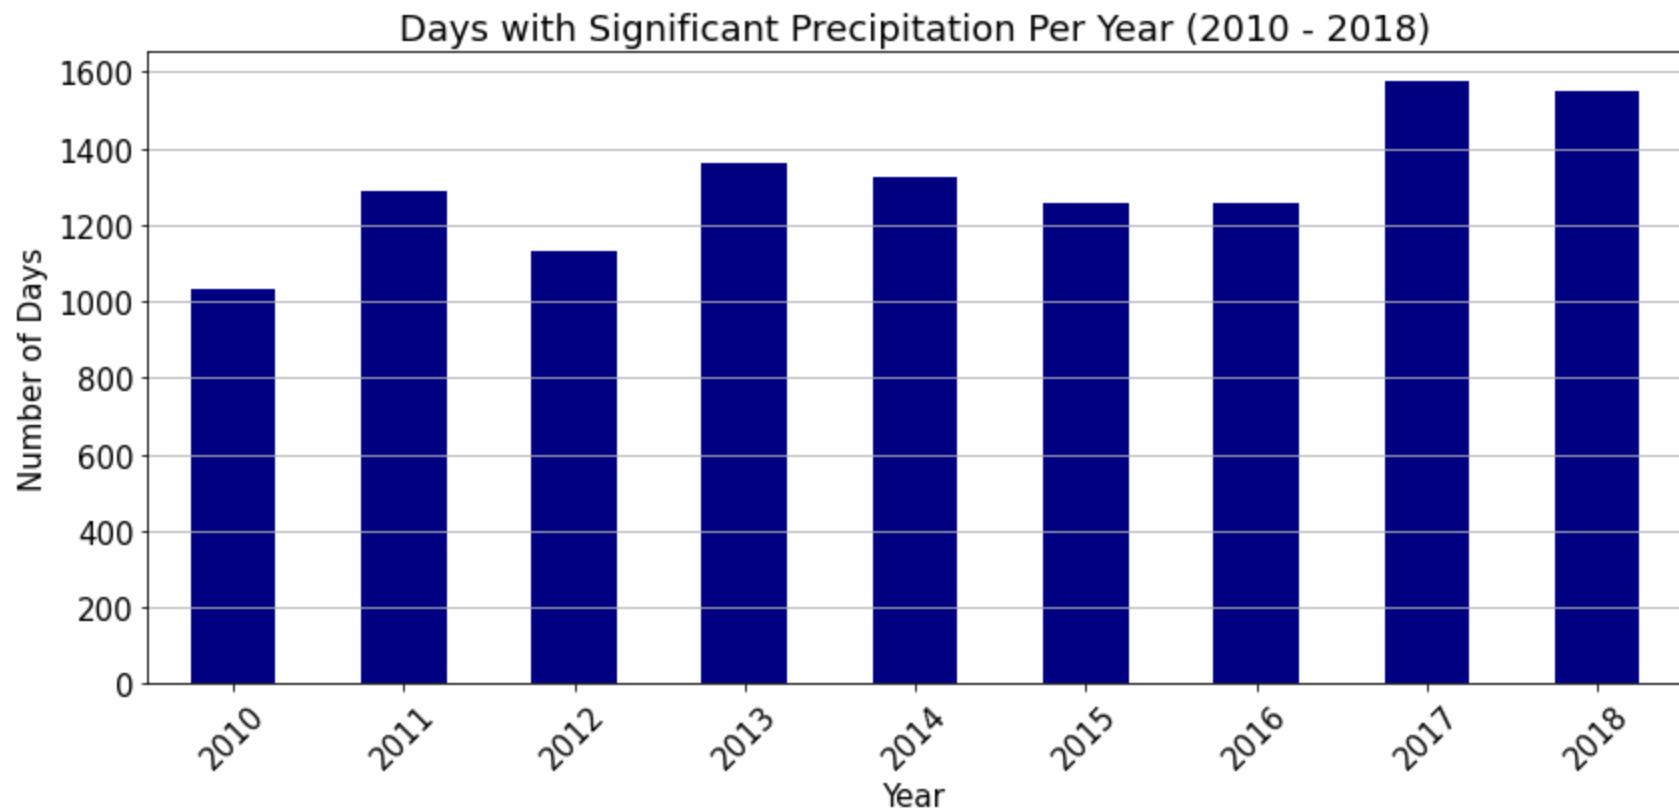
```
# Define a threshold for significant precipitation (in inches)
significant_precipitation_threshold = 0.5

# Filter the dataset for days with precipitation exceeding the threshold
significant_precip_days = weather_data[weather_data['Percipitation'] > significant_precipitation_threshold]

# Count the number of significant precipitation days per year
significant_precip_days_per_year = significant_precip_days.groupby('Year').size()

# Plotting
plt.figure(figsize=(12, 6))
significant_precip_days_per_year.plot(kind='bar', color='navy')
plt.title('Days with Significant Precipitation Per Year (2010 - 2018)')
plt.xlabel('Year')
plt.ylabel('Number of Days')
plt.xticks(rotation=45)
plt.grid(axis='y')

plt.tight_layout()
plt.show()
```



The bar chart illustrates the number of days per year with significant precipitation (more than 0.5 inches) in New York from 2010 through 2018. This visualization helps identify years with higher frequencies of heavy rainfall or significant snowfall events, providing insight into variations in extreme weather conditions over the observed period.

Days with High Wind Speeds

Next, let's define "high wind speeds" as days with wind speeds exceeding 20 mph, which can indicate windy conditions or the presence of storms. We'll count these days per year to observe any trends in their occurrence over time.

In [21]:

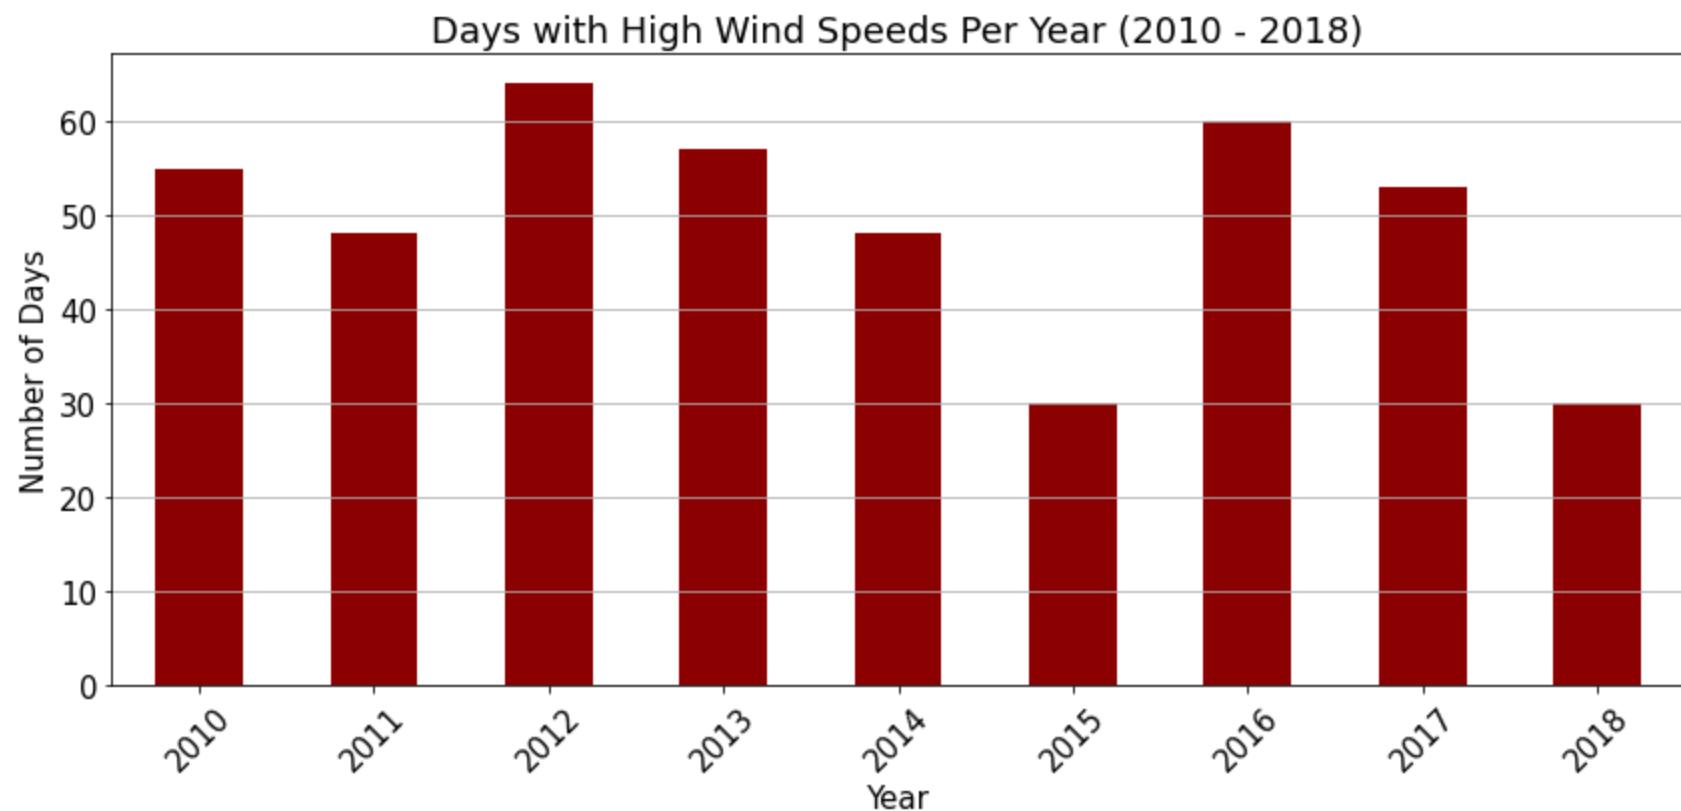
```
# Define a threshold for high wind speeds (in mph)
high_wind_speed_threshold = 20

# Filter the dataset for days with wind speeds exceeding the threshold
high_wind_days = weather_data[weather_data['WindSpeed'] > high_wind_speed_threshold]

# Count the number of high wind speed days per year
high_wind_days_per_year = high_wind_days.groupby('Year').size()
```

```
# Plotting
plt.figure(figsize=(12, 6))
high_wind_days_per_year.plot(kind='bar', color='darkred')
plt.title('Days with High Wind Speeds Per Year (2010 – 2018)')
plt.xlabel('Year')
plt.ylabel('Number of Days')
plt.xticks(rotation=45)
plt.grid(axis='y')

plt.tight_layout()
plt.show()
```



Now that we've taken a quick look at patterns in the two sets of data, we can get into combining the two sets of data, understand correlations and eventually build predictive models to predict number of requests in the next 7 days - as per the original problem description.

First we'll look at monthly trends and then jump into daily trends.

In [22]:

```

requests['Year-Month'] = requests['Created Date'].dt.to_period('M')

# Aggregate 311 requests data to monthly counts
monthly_requests_counts = requests.groupby('Year-Month').size().reset_index(name='Request Count')

# For weather data, ensure we're working with the correct aggregation
weather_data['Year-Month'] = pd.to_datetime(weather_data['Year'].astype(str) + '-' + weather_data['Month'].astype(str))
monthly_weather_agg = weather_data.groupby('Year-Month').agg({
    'MeanTemp': 'mean',
    'Precipitation': 'sum', # Total monthly precipitation
    'WindSpeed': 'mean'
}).reset_index()

# Merging the aggregated monthly 311 requests data with the aggregated monthly weather data
combined_data = pd.merge(monthly_requests_counts, monthly_weather_agg, on='Year-Month', how='inner')

combined_data

```

Out[22]:

	Year-Month	Request Count	MeanTemp	Precipitation	WindSpeed
0	2016-01	27533	28.584615	4957.13	6.978940
1	2016-02	25929	30.937171	5533.10	7.109321
2	2016-03	25768	41.081076	3472.01	6.330385
3	2016-04	24189	44.987842	3785.85	5.931524
4	2016-05	26538	57.836143	4990.16	4.737225
5	2016-06	27325	67.054676	3080.51	5.218597
6	2016-07	26473	73.546323	4815.82	4.880183
7	2016-08	26253	74.225903	4328.30	4.655289
8	2016-09	24831	66.377806	3786.86	4.706922
9	2016-10	26265	53.851048	5964.09	5.575897
10	2016-11	24715	44.245788	5092.11	5.836722
11	2016-12	26889	32.432331	8321.45	6.895344
12	2017-01	27095	32.041296	7810.84	6.382980
13	2017-02	23569	34.722659	5779.87	6.708039

	Year-Month	Request Count	MeanTemp	Precipitation	WindSpeed
14	2017-03	27168	32.805155	4806.76	7.530943
15	2017-04	25453	51.249770	3372.79	6.358111
16	2017-05	27708	56.488759	5394.58	5.947153
17	2017-06	28697	66.863906	6261.59	5.374573
18	2017-07	27487	71.148470	5357.80	4.514880
19	2017-08	26542	68.922875	4826.33	4.363191
20	2017-09	27407	65.132444	3388.15	4.202112
21	2017-10	28675	58.264048	5589.10	5.715326
22	2017-11	28320	41.182797	7389.25	6.249836
23	2017-12	28731	28.129382	7862.71	6.140497
24	2018-01	33165	25.227154	6588.54	6.973660
25	2018-02	25086	33.854462	7815.88	6.044318
26	2018-03	29350	33.732254	7198.82	7.106480
27	2018-04	27842	41.150684	7921.50	6.876334
28	2018-05	30279	62.291871	6491.86	5.966215
29	2018-06	30653	66.491081	5613.94	5.679262
30	2018-07	31126	73.705743	4744.03	5.418868
31	2018-08	31424	73.147894	6476.89	5.233447
32	2018-09	30107	66.046449	4581.64	5.227027
33	2018-10	32840	51.298348	7255.80	6.661797
34	2018-11	30430	44.236401	3092.24	7.176380

As a first pass, let's look at if there's any correlation between precipitation and request count, and wind speed and request count.

In [23]:

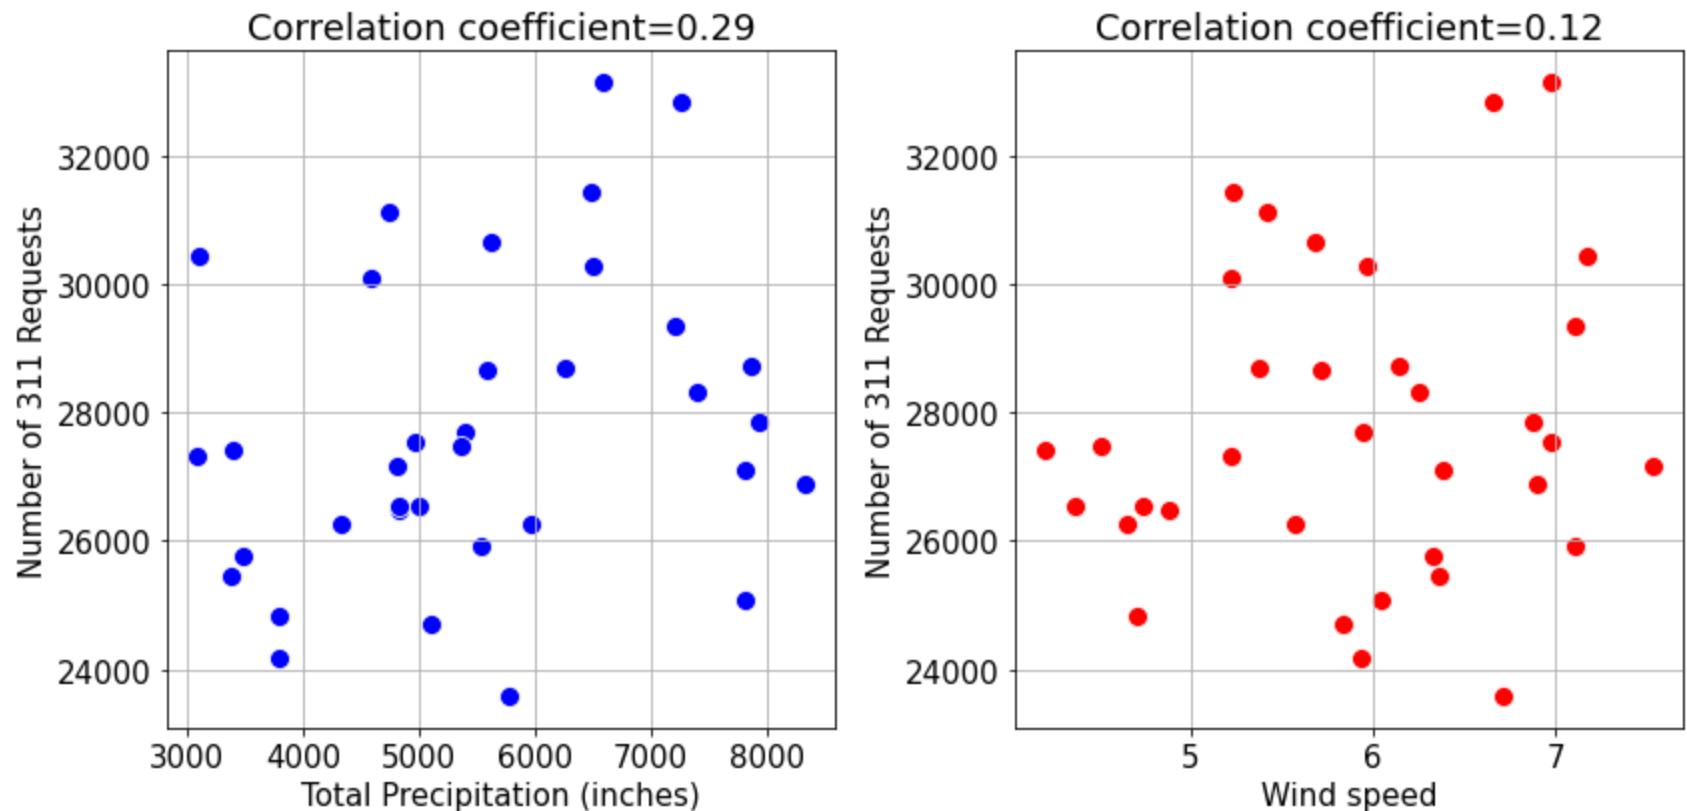
```
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming 'combined_data' is your merged dataset with weather and 311 requests data
fig, axs = plt.subplots(figsize=(12, 6), nrows=1, ncols=2)
```

```
sns.scatterplot(x='Precipitation', y='Request Count', data=combined_data, color='b', s=100, ax=axs[0])
requests_precipitation_corr_coef = combined_data['Request Count'].corr(combined_data['Precipitation'], method='spearman')
axs[0].set_title(f'Correlation coefficient={requests_precipitation_corr_coef:.2f}')
axs[0].set_xlabel('Total Precipitation (inches)')
axs[0].set_ylabel('Number of 311 Requests')
axs[0].grid()

sns.scatterplot(x='WindSpeed', y='Request Count', data=combined_data, color='r', s=100, ax=axs[1])
requests_windspeed_corr_coef = combined_data['Request Count'].corr(combined_data['WindSpeed'], method='spearman')
axs[1].set_title(f'Correlation coefficient={requests_windspeed_corr_coef:.2f}')
axs[1].set_xlabel('Wind speed')
axs[1].set_ylabel('Number of 311 Requests')
axs[1].grid()

plt.tight_layout()
```



We see a weak correlation between precipitation and request count, and almost no correlation between wind speed and request count.

Now, we see some requests with the word heat in them. We can check if there is any correlation between heat-related requests and temperatures.

In [24]:

```
# Filter the 311 dataset for heating complaints
heating_complaints = requests[requests['Complaint Type'].str.contains('Heat', case=False, na=False)]

# Aggregate the count of heating complaints by 'Year-Month'
heating_complaints_monthly = heating_complaints.groupby('Year-Month').size().reset_index(name='Heating Complaints')

# Ensure 'Year-Month' is in a plottable format (if necessary, convert to datetime)
heating_complaints_monthly['Year-Month-Datetime'] = heating_complaints_monthly['Year-Month'].dt.to_timestamp()
heating_complaints_monthly

combined_data_with_heating = pd.merge(monthly_weather_agg, heating_complaints_monthly, on='Year-Month', how='inner')
```

combined_data_with_heating

Out[24]:

	Year-Month	MeanTemp	Precipitation	WindSpeed	Heating Complaints Count	Year-Month-Datetime
0	2016-01	28.584615	4957.13	6.978940	6339	2016-01-01
1	2016-02	30.937171	5533.10	7.109321	5043	2016-02-01
2	2016-03	41.081076	3472.01	6.330385	2223	2016-03-01
3	2016-04	44.987842	3785.85	5.931524	1668	2016-04-01
4	2016-05	57.836143	4990.16	4.737225	1023	2016-05-01
5	2016-06	67.054676	3080.51	5.218597	312	2016-06-01
6	2016-07	73.546323	4815.82	4.880183	290	2016-07-01
7	2016-08	74.225903	4328.30	4.655289	325	2016-08-01
8	2016-09	66.377806	3786.86	4.706922	372	2016-09-01
9	2016-10	53.851048	5964.09	5.575897	2946	2016-10-01
10	2016-11	44.245788	5092.11	5.836722	3702	2016-11-01
11	2016-12	32.432331	8321.45	6.895344	5633	2016-12-01
12	2017-01	32.041296	7810.84	6.382980	4456	2017-01-01
13	2017-02	34.722659	5779.87	6.708039	3126	2017-02-01
14	2017-03	32.805155	4806.76	7.530943	4047	2017-03-01
15	2017-04	51.249770	3372.79	6.358111	1400	2017-04-01
16	2017-05	56.488759	5394.58	5.947153	987	2017-05-01
17	2017-06	66.863906	6261.59	5.374573	443	2017-06-01
18	2017-07	71.148470	5357.80	4.514880	323	2017-07-01
19	2017-08	68.922875	4826.33	4.363191	347	2017-08-01
20	2017-09	65.132444	3388.15	4.202112	358	2017-09-01
21	2017-10	58.264048	5589.10	5.715326	1344	2017-10-01
22	2017-11	41.182797	7389.25	6.249836	4925	2017-11-01
23	2017-12	28.129382	7862.71	6.140497	6388	2017-12-01

	Year-Month	MeanTemp	Precipitation	WindSpeed	Heating Complaints Count	Year-Month-Datetime
24	2018-01	25.227154	6588.54	6.973660	7321	2018-01-01
25	2018-02	33.854462	7815.88	6.044318	2550	2018-02-01
26	2018-03	33.732254	7198.82	7.106480	3186	2018-03-01
27	2018-04	41.150684	7921.50	6.876334	2019	2018-04-01
28	2018-05	62.291871	6491.86	5.966215	657	2018-05-01
29	2018-06	66.491081	5613.94	5.679262	427	2018-06-01
30	2018-07	73.705743	4744.03	5.418868	356	2018-07-01
31	2018-08	73.147894	6476.89	5.233447	351	2018-08-01
32	2018-09	66.046449	4581.64	5.227027	346	2018-09-01
33	2018-10	51.298348	7255.80	6.661797	3547	2018-10-01
34	2018-11	44.236401	3092.24	7.176380	4774	2018-11-01

In [25]:

```

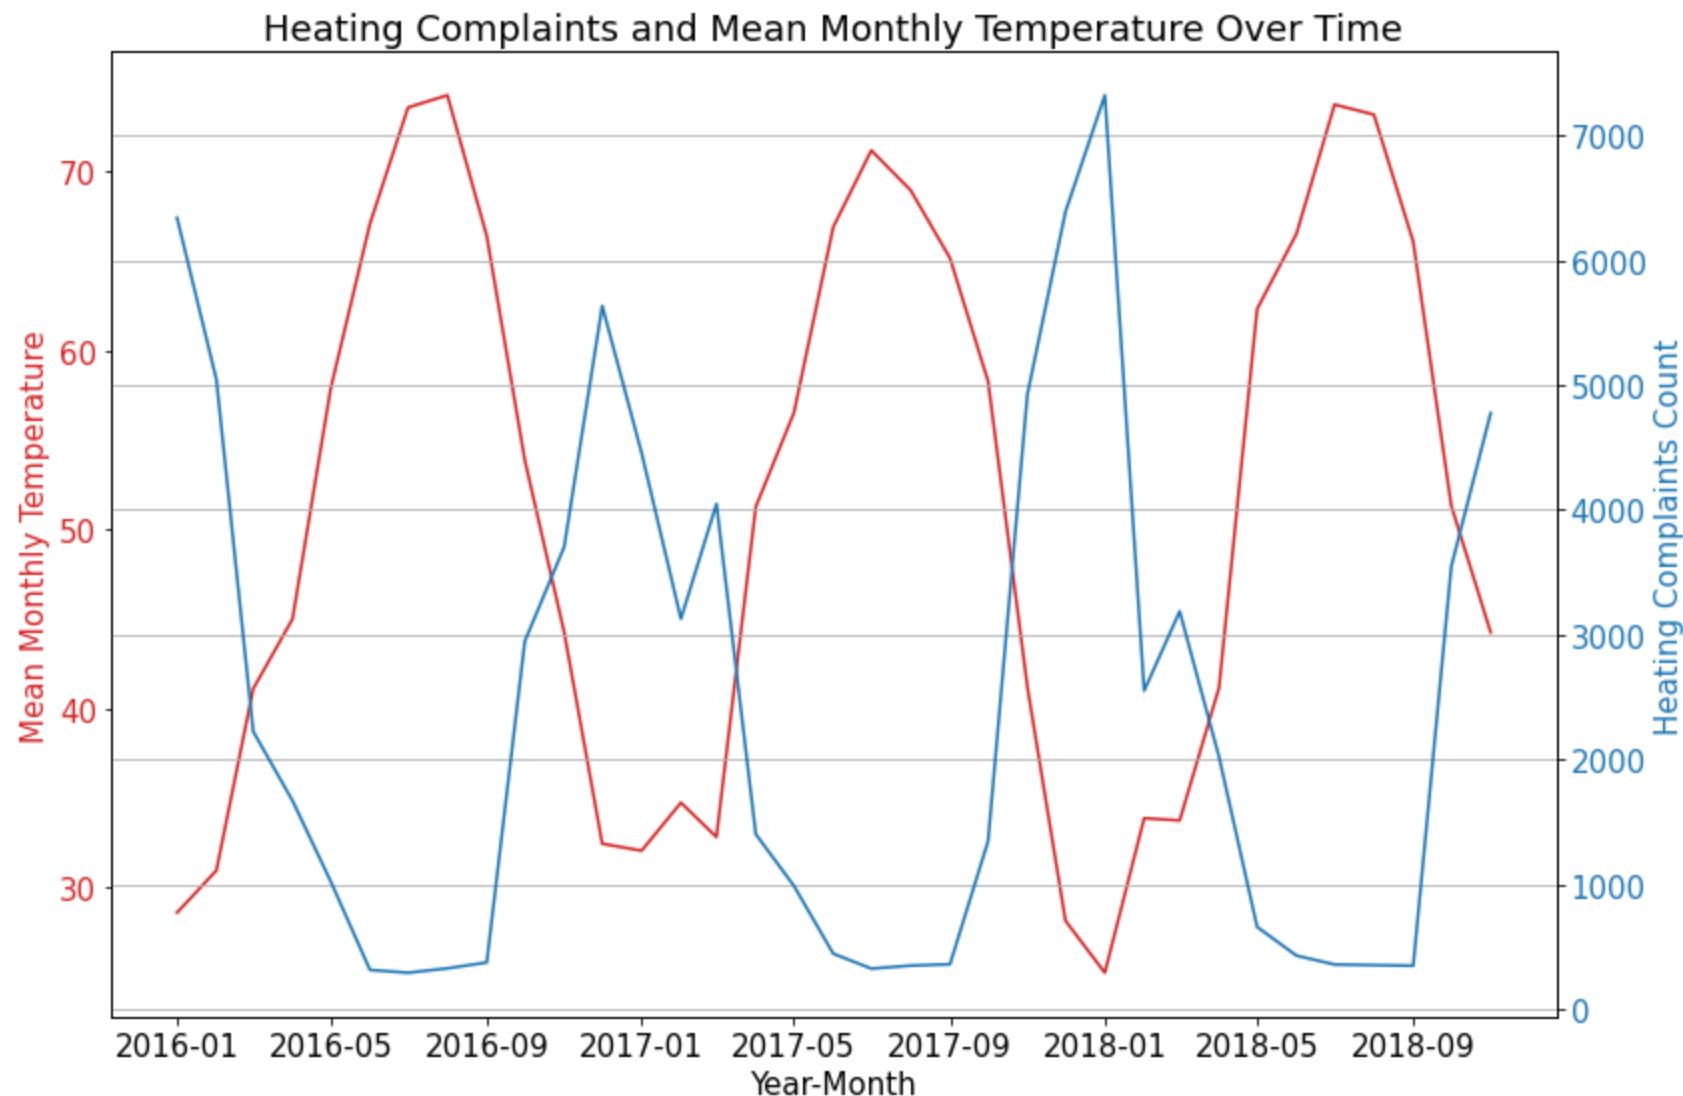
fig, ax1 = plt.subplots(figsize=(12,8))

color = 'tab:red'
ax1.set_xlabel('Year-Month')
ax1.set_ylabel('Mean Monthly Temperature', color=color)
ax1.plot(combined_data_with_heating['Year-Month-Datetime'], combined_data_with_heating['MeanTemp'], color=color)
ax1.tick_params(axis='y', labelcolor=color)

ax2 = ax1.twinx()
color = 'tab:blue'
ax2.set_ylabel('Heating Complaints Count', color=color)
ax2.plot(combined_data_with_heating['Year-Month-Datetime'], combined_data_with_heating['Heating Complaints Count'], color=color)
ax2.tick_params(axis='y', labelcolor=color)

plt.title('Heating Complaints and Mean Monthly Temperature Over Time')
plt.xticks(rotation=45)
fig.tight_layout()
plt.grid()
plt.show()

```



Looks like the heat-related complaints are related to problems with the heaters at home - as these are peaking in winter times. This is unexpected, we can see below that there is a strong negative correlation between the two sets of data.

```
In [26]: combined_data_with_heating['MeanTemp'].corr(combined_data_with_heating['Heating Complaints Count'], method='spe')
Out[26]: -0.930812324929972
```

Now we'll move to daily trends, look at broad correlations and then get into modeling.

In [27]:

```

# Convert 'Created Date' to datetime format, if not already done
requests['Created Date'] = pd.to_datetime(requests['Created Date'])

# Create a new column for the date (without time)
requests['Date'] = requests['Created Date'].dt.date

# Group by the new 'Date' column and count the number of requests per day
daily_requests = requests.groupby('Date').size().reset_index(name='Request Count')

# Assuming 'weather_data' has columns for 'Year', 'Month', and 'Day'
weather_data['Date'] = pd.to_datetime(weather_data[['Year', 'Month', 'Day']]).dt.date

daily_weather_agg = weather_data.groupby('Date').agg({
    'MeanTemp': 'mean', # Daily average mean temperature
    'MaxTemp': 'mean', # Daily average maximum temperature
    'MinTemp': 'mean', # Daily average minimum temperature
    'Precipitation': 'sum', # Total daily precipitation
    'WindSpeed': 'mean', # Daily average wind speed
    'SnowDepth': 'max', # Maximum snow depth for the day
    'Gust': 'max', # Maximum gust speed for the day
    'Rain': 'mean',
    'MaxSustainedWind': 'mean',
    'SnowDepth': 'sum',
    'SnowIce': 'mean'
    # Add any other relevant metrics you wish to include
}).reset_index()

daily_weather_agg['Rain'] = np.round(daily_weather_agg['Rain'])
daily_weather_agg['SnowIce'] = np.round(daily_weather_agg['SnowIce'])

# Merge the daily 311 requests with the daily weather data
combined_data = pd.merge(daily_requests, daily_weather_agg, on='Date', how='inner')

combined_data = combined_data.fillna(0)

combined_data

```

Out[27]:

	Date	Request Count	MeanTemp	MaxTemp	MinTemp	Precipitation	WindSpeed	SnowDepth	Gust	Rain	MaxSustainedWind	Snow
0	2016-01-01	568	34.508333	39.531250	31.058333	302.12	7.821739		4.8	33.0	0.0	15.271795

	Date	Request Count	MeanTemp	MaxTemp	MinTemp	Precipitation	WindSpeed	SnowDepth	Gust	Rain	MaxSustainedWind	SnowIce
1	2016-01-02	610	31.206122	35.873469	27.473469	200.69	7.631915	3.2	31.1	0.0	14.525000	
2	2016-01-03	645	33.993750	39.083333	26.956250	200.30	8.497826	1.2	34.0	0.0	16.974359	
3	2016-01-04	1177	19.645833	35.456250	10.606250	300.59	8.778261	4.4	32.1	0.0	15.941026	
4	2016-01-05	1601	12.697917	26.497917	4.933333	100.10	5.439130	3.2	35.0	0.0	12.584615	
...
1042	2018-11-08	1092	44.509756	54.924390	38.534146	0.18	7.450000	0.0	31.1	0.0	12.907895	
1043	2018-11-09	1020	37.921951	47.680488	31.290244	500.04	5.415789	1.2	27.0	1.0	12.531579	
1044	2018-11-10	867	38.634146	44.358537	29.951220	425.03	11.852632	0.0	42.9	1.0	22.513158	
1045	2018-11-11	814	34.417073	43.041463	28.590244	101.35	7.568421	0.0	34.0	0.0	13.715789	
1046	2018-11-12	941	30.847368	33.242105	28.228947	0.07	2.085714	0.0	0.0	0.0	4.916667	

1047 rows × 12 columns

In [28]:

combined_data.columns

Out[28]:

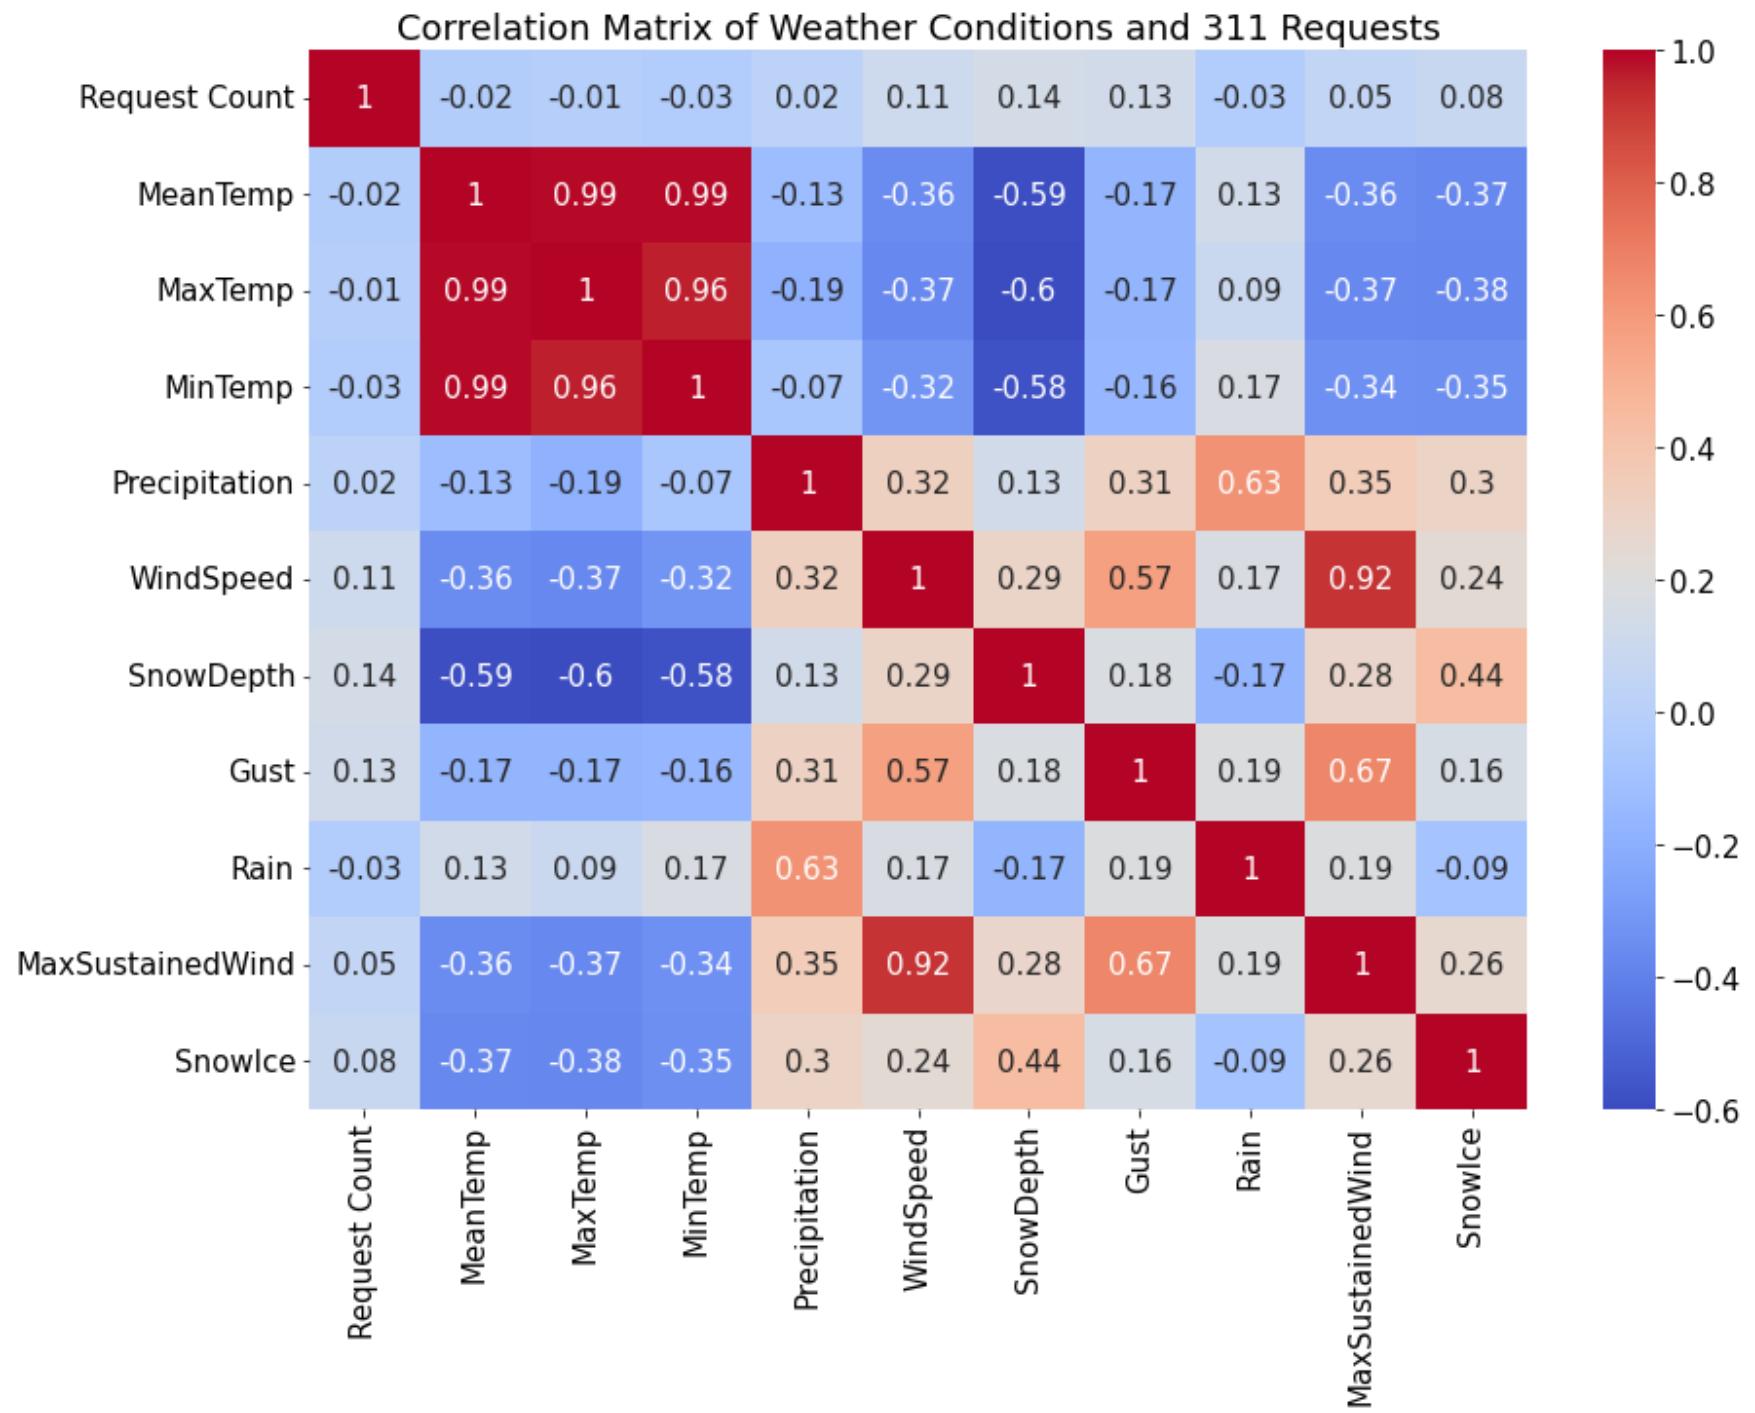
```
Index(['Date', 'Request Count', 'MeanTemp', 'MaxTemp', 'MinTemp',
       'Precipitation', 'WindSpeed', 'SnowDepth', 'Gust', 'Rain',
       'MaxSustainedWind', 'SnowIce'],
      dtype='object')
```

We have included all the available numerical columns in the weather dataset above, and aggregated at the level of each day. Below, we can look at the correlations between these high-level metrics and the corresponding request counts.

In [29]:

```
# Assuming 'combined_data' includes various weather metrics and types of 311 requests
fig, ax = plt.subplots(figsize=(14,10))

correlation_matrix = np.round(combined_data[['Request Count','MeanTemp', 'MaxTemp', 'MinTemp', 'Precipitation',
                                             'MaxSustainedWind', 'SnowIce',]].corr(),2)
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Weather Conditions and 311 Requests')
plt.show()
```



In general, we are seeing low correlations across the board with request count. The most correlated features with request count are `SnowDepth` and `Gust`.

Building the predictive model

Below, we will start working towards the predictive model -

- The eventual goal is to predict the daily 311 inbound calls for the next 7 days.
- We will do this in several steps.
- The last data in the combined dataset is 2018-11-12 - this is where the weather data ends. So, we will first build a quick model to forecast the weather metrics for the next 7 days. To build a quick and reliable forecast, we will use an ARIMA model from python's `statsmodel` library.
- Using these forecasted numbers, we will train a couple of regression models.

In [30]:

```
from statsmodels.tsa.arima.model import ARIMA
# Ignore all warnings
warnings.filterwarnings("ignore")

# Assuming 'weather_data' is your historical DataFrame indexed by datetime
last_date = pd.to_datetime(combined_data.Date.max())
forecast_start_date = last_date + pd.Timedelta(days=1)

# Define the forecast period length (e.g., 7 days)
forecast_period = 7

# Create a date range for the forecast period
forecast_dates = pd.date_range(start=forecast_start_date, periods=forecast_period, freq='D')

def forecast_arima(series, order=(1,1,1), steps=7):
    """
    Forecast the next 'steps' points in the series using ARIMA.

    Parameters:
    - series: pd.Series, the time series data to forecast.
    - order: tuple, the (p, d, q) order of the ARIMA model.
    - steps: int, the number of steps to forecast ahead.

    Returns:
    - forecast: The forecasted values as a pd.Series.
    """
    pass
```

```
model = ARIMA(series, order=order)
model_fit = model.fit()
forecast = model_fit.forecast(steps=steps)
return forecast

weather_metrics = ['MeanTemp', 'MaxTemp', 'MinTemp', 'Percipitation', 'WindSpeed', 'SnowDepth', 'Gust', 'MaxSustWindSpeed']
forecast_results = {}

# Assuming 'weather_data' is your DataFrame and it's indexed by a datetime index
for metric in weather_metrics:
    series = weather_data[metric].dropna() # Drop NA values for simplicity
    forecast = forecast_arima(series, order=(1,1,1), steps=7) # You may need to adjust the order based on the data
    forecast.index = forecast_dates
    forecast_results[metric] = forecast

forecast_results
```

Out[30]:

```
{'MeanTemp': 2018-11-13    55.440374
 2018-11-14    57.587756
 2018-11-15    58.701484
 2018-11-16    59.279113
 2018-11-17    59.578697
 2018-11-18    59.734074
 2018-11-19    59.814660
Freq: D, Name: predicted_mean, dtype: float64,
'MaxTemp': 2018-11-13    67.055139
 2018-11-14    69.345608
 2018-11-15    70.523180
 2018-11-16    71.128591
 2018-11-17    71.439844
 2018-11-18    71.599865
 2018-11-19    71.682135
Freq: D, Name: predicted_mean, dtype: float64,
'MinTemp': 2018-11-13    45.999286
 2018-11-14    48.429294
 2018-11-15    49.522946
 2018-11-16    50.015156
 2018-11-17    50.236681
 2018-11-18    50.336380
 2018-11-19    50.381251
Freq: D, Name: predicted_mean, dtype: float64,
'Percipitation': 2018-11-13    0.090900
 2018-11-14    0.101711
 2018-11-15    0.102996
```

```
2018-11-16    0.103149
2018-11-17    0.103167
2018-11-18    0.103169
2018-11-19    0.103170
Freq: D, Name: predicted_mean, dtype: float64,
'WindSpeed': 2018-11-13    6.438956
2018-11-14    6.608417
2018-11-15    6.639000
2018-11-16    6.644520
2018-11-17    6.645516
2018-11-18    6.645696
2018-11-19    6.645729
Freq: D, Name: predicted_mean, dtype: float64,
'SnowDepth': 2018-11-13    3.000083
2018-11-14    3.673538
2018-11-15    4.127041
2018-11-16    4.432428
2018-11-17    4.638076
2018-11-18    4.776559
2018-11-19    4.869813
Freq: D, Name: predicted_mean, dtype: float64,
'Gust': 2018-11-13    21.590286
2018-11-14    21.815734
2018-11-15    21.847694
2018-11-16    21.852225
2018-11-17    21.852867
2018-11-18    21.852958
2018-11-19    21.852971
Freq: D, Name: predicted_mean, dtype: float64,
'MaxSustainedWind': 2018-11-13    11.490900
2018-11-14    12.019052
2018-11-15    12.126715
2018-11-16    12.148662
2018-11-17    12.153136
2018-11-18    12.154048
2018-11-19    12.154234
Freq: D, Name: predicted_mean, dtype: float64}
```

We now have the weather metrics forecasted for the next 7 days. We'll use this as input in the next few steps to build our predictive model.

Please note the following important considerations during the model building process

- We will first separate the date field into multiple components - such as DayOfWeek, Year, Month, Quarter and Day - this will help us add more granularity into our feature set.
- We will one-hot encode the categorical features - and be sure to use the same attribute set to encode the test set as well (ie, the forecasted weather data).
- When we split the data into train and test sets, we should be mindful to split the data not randomly, but in chronological order, since weather patterns are not random, but are auto-correlated.
- We will measure the performance of the regression models by MAE (mean absolute error) and RMSE (root mean squared error).
- We will compare a simple Linear Regression model with a couple of ensemble models - Random Forest and XGboost.

In [31]:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error

# Extract day of the week as a feature
combined_data['Date'] = pd.to_datetime(combined_data['Date'])
combined_data['DayOfWeek'] = combined_data['Date'].dt.dayofweek
combined_data['Year'] = combined_data['Date'].dt.year
combined_data['Month'] = combined_data['Date'].dt.month
combined_data['Day'] = combined_data['Date'].dt.day
combined_data['Quarter'] = combined_data['Date'].dt.quarter

combined_data = combined_data.sort_values(by='Date')

# Selecting features and target variable for simplicity
features = ['DayOfWeek', 'Year', 'Month', 'Day', 'Quarter', 'MeanTemp', 'MaxTemp', 'MinTemp', 'Percipitation', 'WindSpeed']
target = 'Request Count'

categorical_features = ['DayOfWeek', 'Year', 'Month', 'Day', 'Quarter']

# OneHotEncoding the categorical features
encoded_data = pd.get_dummies(combined_data, columns=categorical_features).drop(columns='Date')

X = encoded_data.drop(columns=target)
y = encoded_data[target]

# Split the data into training and testing sets in chronological order
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, shuffle=False)

# Initialize and train the linear regression model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```

```

# Predict on the testing set
predictions = lr_model.predict(X_test)

# Evaluate the model
mae = mean_absolute_error(y_test, predictions)
rmse = np.sqrt(mean_squared_error(y_test, predictions))

print(f"MAE: {mae}")
print(f"RMSE: {rmse}")

# Assuming forecast_results is a dictionary containing forecasts for each metric
forecast_data = pd.DataFrame(index=forecast_dates)

for metric, forecast in forecast_results.items():
    forecast_data[metric] = forecast

forecast_data = pd.DataFrame(index=forecast_dates)

for metric, forecast in forecast_results.items():
    forecast_data[metric] = forecast

forecast_data['DayOfWeek'] = forecast_data.index.dayofweek
forecast_data['Year'] = forecast_data.index.year
forecast_data['Month'] = forecast_data.index.month
forecast_data['Day'] = forecast_data.index.day
forecast_data['Quarter'] = forecast_data.index.quarter

test_encoded = pd.get_dummies(forecast_data, columns=categorical_features)

test_encoded_aligned = test_encoded.reindex(columns = X.columns, fill_value=0)

request_count_predictions_lr = lr_model.predict(test_encoded_aligned).astype(int)

forecast_lr = pd.DataFrame(request_count_predictions_lr, index=forecast_dates, columns=['Forecast'])

forecast_lr

```

MAE: 88.73617001488095

RMSE: 109.58523427178193

Out[31]:

	Forecast
2018-11-13	1033
2018-11-14	1000
2018-11-15	1000

Forecast

2018-11-16	977
2018-11-17	810
2018-11-18	732
2018-11-19	1007

In [32]:

```
from sklearn.ensemble import RandomForestRegressor

# Initialize the RandomForestRegressor
rf_model = RandomForestRegressor(n_estimators=200, random_state=42, max_depth = 10) # You can adjust these parameters

# Train the model
rf_model.fit(X_train, y_train)

# Predict on the testing set
predictions = rf_model.predict(X_test)

# Evaluate the model
mae = mean_absolute_error(y_test, predictions)
rmse = np.sqrt(mean_squared_error(y_test, predictions))

print(f"MAE: {mae}")
print(f"RMSE: {rmse}")

request_count_predictions_rf = rf_model.predict(test_encoded_aligned).astype(int)
forecast_rf = pd.DataFrame(request_count_predictions_rf, index=forecast_dates, columns=['Forecast'])
forecast_rf
```

MAE: 81.4992087381232

RMSE: 107.48092774985207

Out[32]:

Forecast

2018-11-13	1058
2018-11-14	1060
2018-11-15	1064
2018-11-16	1032
2018-11-17	814

Forecast

2018-11-18	739
2018-11-19	1073

In [33]:

```

import xgboost as xgb

# Initialize the XGBRegressor
xgb_model = xgb.XGBRegressor(objective ='reg:squarederror',
                             colsample_bytree = 0.3,
                             learning_rate = 0.1,
                             max_depth = 5,
                             alpha = 10,
                             n_estimators = 200)

# Train the model
xgb_model.fit(X_train, y_train)

# Predict on the testing set
predictions = xgb_model.predict(X_test)

# Evaluate the model
mae = mean_absolute_error(y_test, predictions)
rmse = np.sqrt(mean_squared_error(y_test, predictions))

print(f"MAE: {mae}")
print(f"RMSE: {rmse}")

request_count_predictions_xgb = xgb_model.predict(test_encoded_aligned).astype(int)
forecast_xgb = pd.DataFrame(request_count_predictions_xgb, index=forecast_dates, columns=['Forecast'])
forecast_xgb

```

MAE: 84.07577136811756

RMSE: 110.15755604624901

Out[33]:

Forecast

2018-11-13	1049
2018-11-14	1011
2018-11-15	1011
2018-11-16	985

Forecast	
2018-11-17	827
2018-11-18	758
2018-11-19	1013

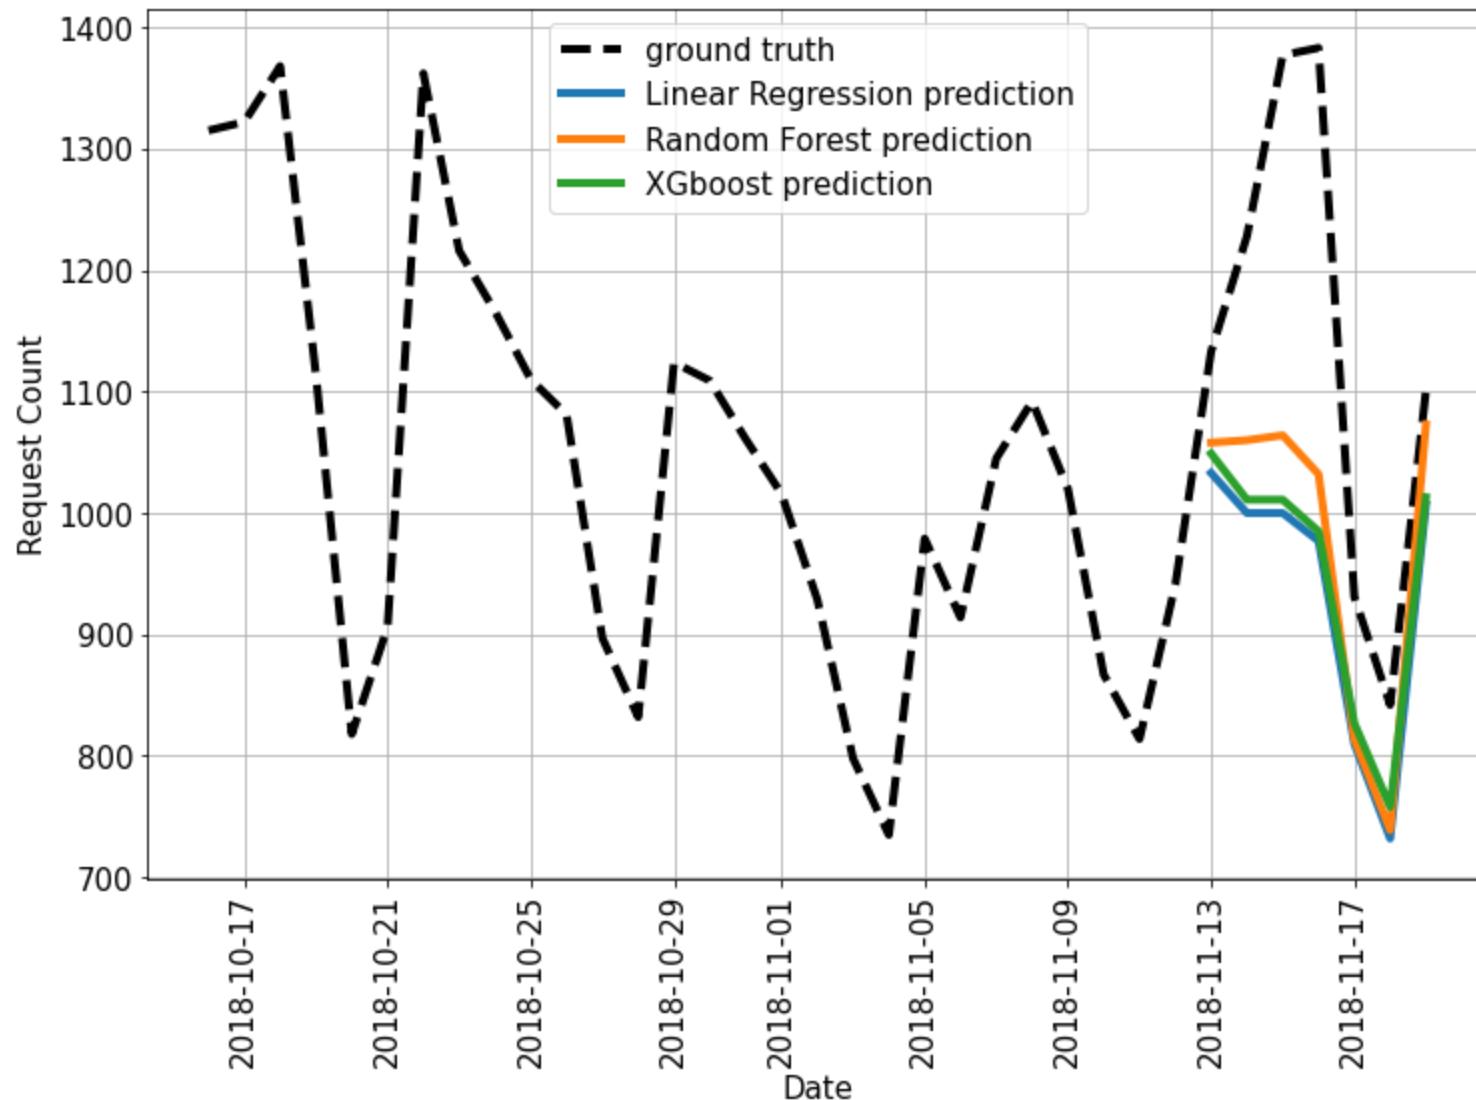
Now, we will visualize the predictions. Since the original dataset also has data for these dates (2018-11-13 through 2018-11-19), we can compare the predictions from the models to this ground truth.

In [34]:

```
fig, ax = plt.subplots(figsize=(12,8))

sns.lineplot(data=daily_requests.loc[(pd.to_datetime(daily_requests['Date']) > pd.to_datetime('2018-10-15')) &
                                         (pd.to_datetime(daily_requests['Date']) <= pd.to_datetime(test_encoded_al:
                                         x='Date', y='Request Count', color='k', label='ground truth', lw=4, linesi
                                         sns.lineplot(x=forecast_lr.index, y=forecast_lr['Forecast'], label='Linear Regression prediction', lw=4)
                                         sns.lineplot(x=forecast_rf.index, y=forecast_rf['Forecast'], label='Random Forest prediction', lw=4)
                                         sns.lineplot(x=forecast_xgb.index, y=forecast_xgb['Forecast'], label='XGboost prediction', lw=4)
                                         plt.setp(ax.get_xticklabels(), rotation=90, ha='center')

                                         ax.legend()
                                         ax.grid()
```



We see that all 3 models performed reasonably well - with Random Forest performing the best. We aren't capturing the large spikes in the request counts (as seen on 2018-11-16), but otherwise, our predictions aren't too bad.

Summary and Conclusion

We analyzed the 311 requests data from NY, looked at various attributes and patterns in the data. While there are numerous types of complaints/requests in the dataset, a subset of these requests could be influenced by weather patterns. We eventually looked at the

weather dataset for the same time period, and then used the weather features to predict the number of daily requests for the next 7 days. We do this by joining the two datasets, creating multiple features from the available data to inform different regression models. Following are the main takeaways

- the input dataset of requests was very large - we tackle this by randomly sampling 1M rows of data spanning the entire time period.
- we corrected for erroneous values in the data, such as negative resolution times.
- given more time, we could have dedicated more efforts towards digging deeper into the datasets to understand more nuanced patterns. For instance, a better understanding about the different types of complaints would enable us to develop a conceptual understanding about which types of requests are affected by weather.
- along the same lines, we could have also used other publicly available datasets to improve our predictions. For instance, one of the most frequent requests was "Noise - Residential"; one could expect some correlations between the prevalence of this request and proximity to urban areas, since one would residential noise complaints to be more prevalent in apartment buildings rather than single-family houses. And this applies for other types of noise complaints as well.
- there are a number of other types of complaints which are related to living conditions - such as unsanitary conditions, construction-related issues etc, where we can again benefit from housing-quality related data.
- we could also expect improved accuracies by running the above models at more specific locations - ie, using local weather patterns to predict the request count for the borough or some approximated lat/long combination. This is possible since we have lat/long information in both datasets.
- finally, given more time, we could've continued to improve the model parameters - through hyperparameter tuning, and eventually getting into more complex models. Overall, this was a fun exercise.

In []: